

Jonny L. Saunders, Lucas A. Ott, Michael Wehr@

University of Oregon

Institute of Neuroscience, Department of Psychology

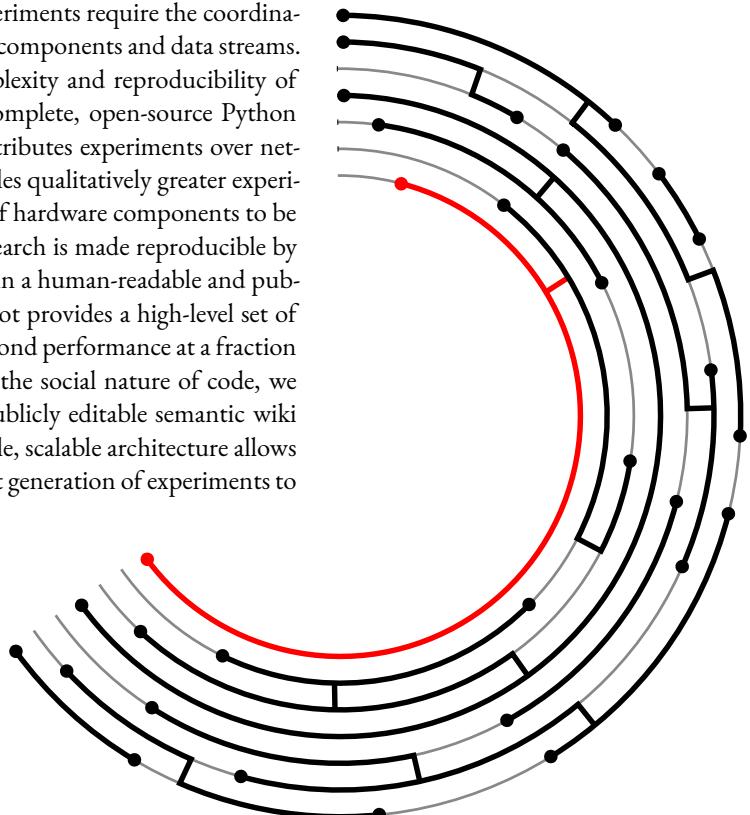
Eugene, OR 97403, United States

AUTOPilot

Automating experiments with lots of Raspberry Pis

Neuroscience needs behavior, and behavioral experiments require the coordination of large numbers of heterogeneous hardware components and data streams. Currently available tools strongly limit the complexity and reproducibility of experiments. Here we introduce Autopilot, a complete, open-source Python framework for experimental automation that distributes experiments over networked swarms of Raspberry Pis. Autopilot enables qualitatively greater experimental flexibility by allowing arbitrary numbers of hardware components to be combined in arbitrary experimental designs. Research is made reproducible by documenting all data and task design parameters in a human-readable and publishable format at the time of collection. Autopilot provides a high-level set of programming tools while maintaining submillisecond performance at a fraction of the cost of traditional tools. Taking seriously the social nature of code, we scaffold shared knowledge and practice with a publicly editable semantic wiki and a permissive plugin system. Autopilot's flexible, scalable architecture allows neuroscientists to work together to design the next generation of experiments to investigate the behaving brain.

version: v2-biorxiv-8-g70c8854



 DOCS

 SOURCE

 WIKI

 PAPER SOURCE

 PAPER PLUGIN

We would like to acknowledge and thank Lucas Ott and Tillie Morris for doing most of the behavioral training and being so patient with the bugs, Brynna Paros and Nick Sattler for their help with constructing our behavioral boxes, Chris Rogers who has been brave enough to adopt and contribute to Autopilot in its roughest state, Arne Meyer, Mikkel Roald-Arbøl, and David Robbe who have contributed code and advice, Mackenzie Mathis, Alex Mathis, Gonçalo Lopes, and Gary Kane who collaborated on the DeepLabCut-Live project and provided many a mentorship along the way, Jeremy Delahanty for his inspiring tenacity and bumbleness in thinking about better research tools, Matt Smear and Reese Findley for loaning us their Bpod for far longer than they intended to, John Boosinger and the rest of the staff in the machine shop for all their advice and letting me use all their tools, Erik Flister whose Ratrix software inspired some of the design features of Autopilot [1], Santiago Jaramillo whose [TASKontrol](#)[2] gave inspiration for GUI design and served as an early scaffolding to learn Python, my labmates Molly Shallow and Sam Mehan who kept me afloat in my last months of dissertation writing, Rocky Penick for her help strapping Autopilot onto the back of Evan Vickers' mesoscope rig (and Evan for letting us play with his rig), several artists on [flaticon.com](#) ([Freepik](#), [Nikita Golubev](#), [Those Icons](#)) whose work served as stems for some of the figures, and the Janet Smith House for the endless support and relentless criticism of the figures. This material is based on work supported by NIH NIDCD R01 DC-015828, NSF Graduate Research Fellowship No. 1309047, and a University of Oregon Incubating Interdisciplinary Initiatives award.

Contribution Statement: JLS designed and wrote the software, documentation, wiki, figures, ran the tests, and wrote and edited the paper. LO trained the animals and did an extensive amount of beta testing, bugfinding, and made some of the hardware designs on the wiki. MW mentored, edited the paper, and beta tested the software.

Contents

1	Introduction	5
1.1	Existing Systems for Behavioral Experiments	7
1.2	Limitations of Existing Systems	8
2	Design	11
2.1	Flexibility	11
2.2	Reproducibility	13
3	Methods & Results	17
3.1	Design	17
3.2	Data	19
3.3	Tasks	21
3.4	Hardware	26
3.5	Transforms	27
3.6	Stimuli	28
3.7	Agents	28
3.8	Networking	30
3.9	GUI & Plots	32
4	Tests	33
4.1	GPIO Latency	33
4.2	Sound Latency	35
4.3	Network Latency	35
4.4	Network Bandwidth	36
4.5	Distributed Go/No-go Task	38
5	Limitations and Future Directions	41
6	Glossary	43
	Bibliography	45

1

Introduction

1 ANIMAL BEHAVIOR experiments need precision and patience, so we make computers do them for us. The complexity of contemporary behavioral experiments,
2 however, presents a stiff methodological challenge. For example, researchers might
3 wish to measure pupil dilation[3], respiration[4], and running speed[5], while tracking
4 the positions of body parts in 3 dimensions[6] and recording the activity of large
5 ensembles of neurons[7], as subjects perform tasks with custom input devices such
6 as a steering wheel[8] while immersed in virtual reality environments using stimuli
7 synthesized in real time[9, 10]. Coordinating the array of necessary hardware into
8 a coherent experimental design—with the millisecond precision required to study
9 the brain—can be daunting.

11 Historically, researchers have developed software to automate behavior experiments
12 as-needed within their lab or relied on purchasing proprietary software (eg. [11]).
13 Open-source alternatives have emerged recently, often developed in tandem with
14 hardware peripherals available for purchase [12, 13]. However, the diverse hard-
15 ware and software requirements for behavioral experiments often lead researchers
16 to cobble together multiple tools to perform even moderately complex experiments.
17 Understandably, most software packages do not attempt to simultaneously support
18 custom hardware operation, behavioral task logic, stimulus generation, and data
19 acquisition. The difficulty of designing and maintaining lab-idiosyncratic systems
20 thus defines much of the everyday practice of science. Idiosyncratic systems can hin-
21 der reproducibility, especially if the level of detail reported in a methods section is
22 sparse[14]. Additionally, development time and proprietary software are expensive,
23 as are the custom hardware peripherals that are required to use most available open-
24 source behavior software, stratifying access to state-of-the-art techniques according
25 to inequitable funding distributions.

26 Technical challenges are never merely technical: they reflect and are structured by
27 underlying *social* challenges in the organization of scientific labor and knowledge
28 work. Lab infrastructure occupies a space between technology intended for indi-
29 vidual users and for large organizations: that of *groupware*¹ [16, 15]. Experimental
30 frameworks thus face the joint challenge of technical competency while also embed-
31 ding in and supporting existing cultures of practice. Behind every line of code is
32 an unwritten wealth of technical knowledge needed to make use of it, as well as an
33 unspoken set of beliefs about how it is to be used — labs aren’t born fresh on re-
34 lease day ready to retool at a moment’s notice, they’re held together by decades of
35 duct tape and run on ritual. The boundaries of this “contextual knowledge” extend
36 fluidly beyond individual labs, structuring disciplinary, status, and role systems in
37 scientific work[17]. Given their position at the intersection of scientific theory, tech-
38 nical work, data production, and social organization, experimental frameworks are
39 an elusive design challenge, but also an underexplored means of realizing some of
40 our loftier dreams of open, accessible, and collaborative science.

¹ “Our original definition of groupware was ‘intentional group processes plus software to support them.’ It has both *computer* and *human* components: software of the computer and ‘software’ of the people using it. [...] Recently this definition has been extended to include other more expressly cultural factors including myth, values and norms. The computer software should reflect and support a group’s purpose, process and culture.”

Peter and Trudy Johnson-Lenz (1991)[15]

⁴¹ Here we present Autopilot, a complete open-source software and hardware frame-
⁴² work for behavioral experiments. We leverage the power of distributed computing
⁴³ using the surprisingly capable Raspberry Pi 4² to allow researchers to coordinate ar-
⁴⁴bitrary numbers of heterogeneous hardware components in arbitrary experimental
⁴⁵ designs.

⁴⁶ Autopilot takes a different approach than existing systems to overcome the technical
⁴⁷ challenges of behavioral research: *just use more computers*. Specifically, the advent of
⁴⁸ inexpensive single-board computers (ie. the Raspberry Pi) that are powerful enough
⁴⁹ to run a full Linux operating system allows a unified platform to run on every Pi
⁵⁰ or other computer in the system so that they can work together seamlessly. At the
⁵¹ core of its architecture are networking classes (Section 3.8) that are fast enough to
⁵² stream electrophysiological or imaging data and flexible enough to make the mutual
⁵³ coordination of hardware straightforward.

⁵⁴ This distributed design also makes Autopilot extremely scalable, as the Raspberry
⁵⁵ Pi's \$35-\$75 price tag makes it an order of magnitude less costly than comparable
⁵⁶ systems (Section 2.2). Its low cost doesn't come at the expense of performance or
⁵⁷ usability: Autopilot provides an approachable, high-level set of tools that still have
⁵⁸ input and output precision between dozens of microseconds to a few milliseconds
⁵⁹ (Sections ?? and 4).

⁶⁰ Autopilot balances experimental flexibility with support. Its task design infrastruc-
⁶¹ture is flexible enough to perform arbitrary experiments, but also provides support
⁶² for data management, plotting task progress, and custom training regimens. We
⁶³ try to bridge multiple modalities of use: use its modular framework of tools out of
⁶⁴ the box, or use its [complete low-level API documentation](#)³ to hack it to do what
⁶⁵ you need. Rather than relying on costly proprietary hardware modules, users can
⁶⁶ take advantage of the wide array of peripherals and extensive community support
⁶⁷ available for the Raspberry Pi. Autopilot is designed to be *permissive*: build your
⁶⁸ whole experiment with it or just use its networking modules, adapt it to existing
⁶⁹ hardware, integrate your favorite analysis tool. We designed Autopilot to *play nice*
⁷⁰ with other software libraries and existing practices rather than force you to retool
⁷¹ your lab around it.

⁷² Finally, we have designed Autopilot to help scientists do reproducible research and
⁷³ be good stewards of the human knowledge project. Experiments are not written as
⁷⁴ scripts that are reliant on the particularities of each researcher's hardware configu-
⁷⁵ration. Instead, we have designed the system to encourage users to write reusable,
⁷⁶ portable experiments that can be incorporated into a public central library while
⁷⁷ also allowing space to iterate and refine without needing to learn complicated pro-
⁷⁸gramming best-practices to contribute. Every parameter that defines an experiment
⁷⁹ is automatically saved in publication-ready format, removing ambiguity in reported
⁸⁰ methods and facilitating exact replication with a single file. Its plugin system is built
⁸¹ atop a densely-linked [semantic wiki](#)⁴ that fluidly combines human- and computer-
⁸² readable, communally editable technical knowledge that surrounds your experiments
⁸³ with the software that performs them.

⁸⁴ We begin by defining the requirements of a complete behavioral system and evalua-
⁸⁵ting two current examples (Sections 1.1 and 1.2). We then describe Autopilot's
⁸⁶design principles (Section 2) and how they are implemented in the program's struc-
⁸⁷ture (Section 3). We close with a demonstration of its current capabilities and our
⁸⁸plans to expand them (Sections 4 and 5).

² See Table 3.2

³ <https://docs.auto-pi-lot.com>

⁴ <https://wiki.auto-pi-lot.com>

89 *1.1 Existing Systems for Behavioral Experiments*

90 At minimum, a complete system to automate behavioral experiments has 6 require-
91 ments:

- 92 1. **Hardware** to interact with the experimental subject, including **sensors** (eg. pho-
93 todiodes, cameras, rotary encoders) to receive input and **actuators** (eg. lights,
94 motors, solenoids) to provide feedback.
- 95 2. Some capability to synthesize and present sensory **stimuli**. Ideally both discrete
96 stimuli, like individual tone pips or grating patches, and continuous stimuli, like
97 those used in virtual reality experiments, should be possible.
- 98 3. A framework to coordinate hardware and stimuli as a **task**. Task definition
99 should be flexible such that it facilitates rather than constrains experimental de-
100 sign.
- 101 4. A **data management** system that allows fine control of data collection and format.
102 Data should be human readable and include complete metadata that allows inde-
103 pendent analysis and reproduction. Ideally the program would also allow some
104 means of realtime data processing of sensor values for use in a task.
- 105 5. Some means of **visualizing data** as it is collected in order to observe task status. It
106 should be possible to customize visualization to the needs and structure of the
107 task.
- 108 6. Finally, a **user interface** to control task operation. The UI should make it possible
109 for someone who does not program to operate the system.

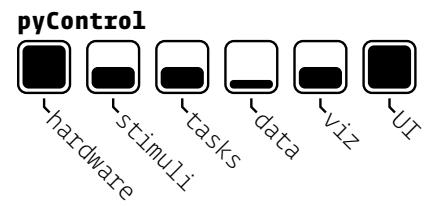
110 We will briefly describe two other systems that meet this definition of completeness:
111 pyControl and Bpod.

112 *pyControl*

113 pyControl[18] is a behavioral framework built in Python by the Champalimaud
114 Foundation. It uses the **micropython microcontroller** (“pyboard”) as its primary
115 hardware device along with several extension boards **sold by openephys**. The py-
116 board has four I/O ports, or eight with a multiplexing expander board. Schematics
117 are available for many other hardware components like solenoid valve drivers and
118 rotary encoders. Multiple pyboards can be connected to a computer via USB and
119 run independent tasks simultaneously with a GUI.

120 There is limited support for some parametrically defined sound stimuli, presented
121 from a separate amplifier connected using the I2C protocol. Visual stimuli are un-
122 supported.

123 Like most behavioral software, pyControl uses a finite-state machine formalism to
124 define its tasks. A task is a set of discrete states, each of which has a set of events that
125 transition the task from one state to another. pyControl also allows timed transi-
126 tions between states, and one function that is called on every event for a rough sort
127 of parallelism. pyControl also allows the use of external variables to control state
128 logic, making these state machines more flexible than strict finite state machines.



D 0 2
D 8976 3
D 8976 1
P 8976 Print Statement
D 10162 3
D 10163 2

Figure 1.1: pyControl data is stored as plain text, each line having a type (like **Data** or **Print**), timestamp, and state

129 All events and states are stored alongside timestamps as a plain text log file, one file
 130 per subject per session (Figure 1.1). Analog data are stored in a custom binary ser-
 131 alization that alternates 4-byte data and timestamp integers.

132 There is only one plot type available in the GUI, a raster plot of events, and no facil-
 133 ity for varying the plot by task type. The GUI is otherwise quite capable, including
 134 the ability to batch run subjects, redefine task variables, and configure hardware.

135

136 *Bpod*

137 Bpod is primarily a collection of hardware designs and an assembly service run by
 138 [Sanworks LLC](#). Similar to pyControl, each Bpod behavior box is based on a finite-
 139 state machine microcontroller with four I/O ports. Additional hardware modules
 140 provide extended functionality. Bpod is controlled using its own [MATLAB pack-](#)
 141 [age](#), though there are at least two other third-party software packages, [BControl](#) and
 142 [pyBpod](#), that can control Bpod hardware. A task is implemented as a MATLAB
 143 script that constructs a new state machine for each trial, uploads it to the Bpod, and
 144 waits for the trial to finish. As a result, only one Bpod can be used per host computer,
 145 or at least per MATLAB session. Data are stored as trial-split events in a MATLAB
 146 structure.

147 There are a few basic plots for two-alternative forced choice tasks, but there doesn't
 148 seem to be a prescribed way to add additional plots. Bpod has a reasonably complete
 149 GUI for managing the hardware and running tasks, but it is relatively technical (Fig-
 150 ure 1.2).

151 For brevity we have omitted many other excellent tools that perform some subset
 152 of the operations of a complete behavioral system, or otherwise have a substantial
 153 difference in scope.⁵

154 1.2 Limitations of Existing Systems

155 We see several limitations with these and other behavioral systems:

- 156 • **Hardware** - Both Pycontrol and Bpod strongly encourage users to purchase a
 157 limited set of hardware modules and add-ons from their particular hardware
 158 ecosystem. If a required part is not available for purchase, neither system pro-
 159 vides a clear means of interacting with custom hardware aside from typical digi-
 160 tal inputs and outputs, requiring the user to ‘tack on’ loosely-integrated compo-
 161 nents. There is also a hard limit on the *number* of hardware peripherals that can
 162 be used in any given task, as there is no ability to use additional pyboards or Bpod
 163 state machines in a single task. The microcontrollers used in these systems also
 164 impose strong limits on their software: neither run a full, high-level program-
 165 ming language⁶. We will discuss this further in section 2.1. A broader limitation
 166 of existing systems is the difficulty of flexibly integrating diverse hardware with
 167 the analytical tools necessary to perform the next generation of behavioral neuro-
 168 science experiments that study “naturalistic, unrestrained, and minimally shaped
 169 behavior”[26].

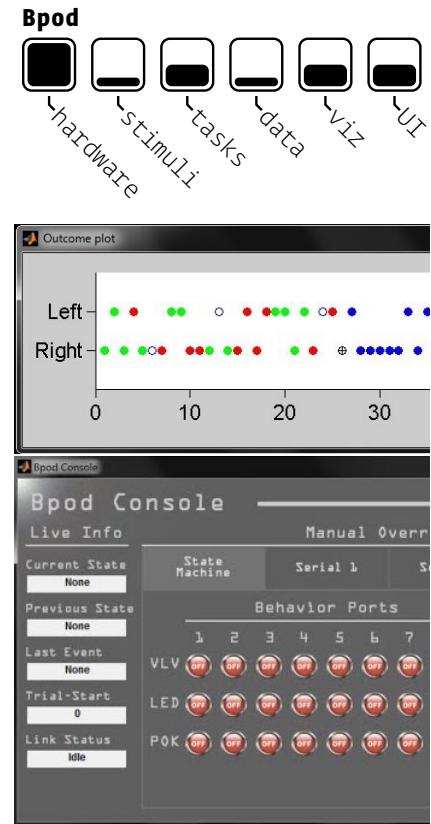


Figure 1.2: A Bpod event plot (above) showing the results of individual behavioral trials, and the Bpod GUI (below).

⁵ Other tools:

- Bonsai[19] - [site](#), [git](#)
- Expyriment[20] - [site](#), [git](#)
- PsychoPy[21] - [site](#), [git](#)
- OpenSesame[22] - [site](#), [git](#)
- SMiLE - [docs](#)
- ArControl[23] - [git](#)
- and see [OpenBehavior](#)

⁶ Bpod runs [custom firmware](#) written in C++ on a [Teensy 3.6](#) microcontroller. pyControl's pyboard runs [micropython](#), a subset of Python that excludes canonical libraries like [numpy](#)[24] or [scipy](#)[25]

170 • **Stimuli** - Stimuli are not tightly integrated into either of these systems, requiring
 171 the user to write custom routines for their synthesis, presentation, and descrip-
 172 tion in the resulting data. Neither are capable of delivering visual stimuli. Since
 173 the publication of the initial version of this manuscript, Bpod has added sup-
 174 port for a HiFiBerry sound card that we also describe here[27], but the sound
 175 generation API appears to be [unchanged](#), with a single method for generating
 176 [sine waves](#). Some parametric audio stimuli are included in the [pyControl source](#)
 177 [code](#) but we were unable to find any documentation or examples of their use.

178 • **Tasks** - Tasks in both systems require a large amount of code and effort dupli-
 179 cation. Neither system has a notion of reusable tasks or task ‘templates,’ so every
 180 user typically needs to rewrite every task from scratch. Bpod’s structure in par-
 181 ticular tends to encourage users to write long [task scripts](#) that contain the entire
 182 logic of the task including updating plots and recreating state machines (Figure
 183 [1.3](#)). Since there is little notion of how to share and reuse common operations,
 184 most users end up creating their own secondary libraries and writing them from
 185 scratch. Another factor that contributes to the difficulty of task design in these
 186 systems is the need to work around the limitations of finite state machines, which
 187 we discuss further in section [3.3](#).

188 • **Data** - Data storage and formatting is basic, requiring extensive additional pro-
 189 cessing to make it human readable. For example, to determine whether a subject
 190 got a trial correct in an [example](#) Bpod experiment, one would use the following
 191 code:

```
192 SessionData.RawEvents.Trial{1,1}.States.Punish(1) ~= NaN
```

193 As a result, data format is idiosyncratic to each user, making data sharing depen-
 194 dent on manual annotation and metadata curation from investigators.

195 • **Visualization & GUI** - The GUIs of each of these systems are highly technical,
 196 and are not designed to be easily used by non-programmers, though pyControl’s
 197 documentation offsets much of this difficulty. Visualization of task progress is
 198 quite rigid in both systems, either a timeseries of task states or plots specific to
 199 two-alternative forced choice tasks. In the examples we have seen, adapting plots
 200 to specific tasks is mostly ad-hoc use of external tools.

201 • **Documentation** - Writing good documentation is challenging, but particularly
 202 for infrastructural systems where a user is likely to need to modify it to suit their
 203 needs it is important that it be possible to understand its lower-level workings.
 204 PyControl has relatively good [user documentation](#) for how to use the system,
 205 but no API-level documentation. Bpod’s [documentation](#) is a bit more scattered,
 206 and though it does have documentation for a [subset of its functions](#), there is little
 207 indication of how they work together or how someone might be able to modify
 208 them.

209 • **Reproducibility** - As of [November 2020](#), pyControl has [versioned task files](#)
 210 that append a hash to each version of a task and save it along with any produced
 211 data, tying the data to exactly the code that produced it. PyControl’s most recent
 212 releases have explicit [version numbers](#), but these don’t appear to be saved along
 213 with the data. Bpod stores neither code nor task versions in its data. Neither
 214 system saves experimental parameter changes by default—and the GUIs of both
 215 allow parameters to be changed at will—and so critical data could be lost and ex-
 216 periments made unreplicable unless the user writes custom code to save them.

```
for currentTrial = 1:MaxTrials
% new state machine every trial
sma = NewStateMachine();

% add states and transitions
sma = AddState(sma,
  'Name', 'Wait', ...
  'Timer', 0, ...
  'StateChangeConditions', ...
  {'Port2In', 'Delay'}, ...
  'OutputActions', ...
  {'AudioPlayer1','*'});

% add more states ...

% upload and run task
SendStateMatrix(sma);
RawEvents = RunStateMatrix;

% manually gather data and params
BpodSystem.Data = AddTrialEvents(
  BpodSystem.Data, RawEvents);

% plotting in the main loop
UpdateSideOutcomePlot( ... );
UpdateTotalRewardDisplay( ... );

% manually save data
SaveBpodSessionData;
end
```

Figure 1.3: Bpod’s general task structure.

217 Bpod has an undocumented [plugin system](#), but neither system has a formal sys-
 218 tem for sharing plugins or task code, requiring work to be duplicated across all
 219 users of the system.

- 220 • **Integration and Extension** - Integration with other systems that might han-
 221 dle some out-of-scope function is tricky in both of these example systems. All
 222 systems have some limitation, so care must be taken to provide points by which
 223 other systems might interact with them. One particularly potent example is the
 224 use of Bpod in the International Brain Laboratory's standardized experimental
 225 rig[28], which relies on a single-purpose [93 page PDF](#) to describe how to use
 226 the [iblrig](#) library, which consists of a large amount of single-purpose code for
 227 stitching together pybpod with [bonsai](#) for controlling video acquisition. Even
 228 if a system takes a large amount of additional work to integrate with another,
 229 hopefully the system allows it to be done in a way such that it can be reused and
 230 shared with others in the future so they can be spared the trouble. The relatively
 231 sparse [documentation](#) and the high proportion of ibl-specific code present in the
 232 repository make that seem unlikely.

233 Some of these limitations are cosmetic—fixable with additional code or hardware—
 234 but several of the most crucial are intrinsic to the design of these systems.

235 These systems, among others, have pioneered the development of modern behav-
 236 ioral hardware and software, and are to be commended for being open-source and
 237 highly functional. One need look no further for evidence of their usefulness than
 238 to their adoption by many labs worldwide. At the time that these systems were de-
 239 veloped, a general-purpose single-board computer with performance like the Rasp-
 240 berry Pi 4 was not widely available. The above two systems are not unique in their
 241 limitations⁷, but are reflective of broader constraints of developing experimental
 242 tools: solving these problems is *hard*. We are only able to articulate the design prin-
 243 ciples that differentiate Autopilot by building on their work.

⁷ And Autopilot, of course, also has many of its own weaknesses

244 2

245 *Design*

246 2.1 *Flexibility*

247 *Single-language*

248 Behavior software that uses dedicated microprocessors must have some routine for
249 compiling the high-level abstraction of the experiment into machine code. This
250 gives those systems a theoretical advantage in processing speed, but the compiler
251 becomes the bottleneck of complexity: only those things that can be compiled can
252 be included in the experiment. This may in part contribute to the ubiquity of state-
253 machine formalisms in behavior software.

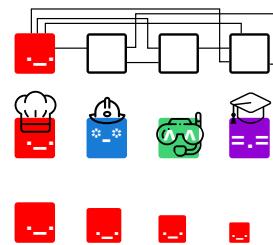
254 Because Python is used throughout the system, extending Autopilot's functionality
255 is straightforward. Task design (see section 3.3) is effectively arbitrary—anything
256 that can be expressed in Python is a valid task. This also allows Autopilot to easily
257 be extended to make use of external libraries (eg. our integration with DeepLabCut-
258 Live[29] and our **planned** integration with OpenEphys).

259 *Modularity*

260 Although Autopilot deeply integrates with the Raspberry Pi's hardware, we have
261 also worked to make its components modular. There is a tension between providing
262 a full-featured behavioral system and the flexibility of its components — as additional
263 features are added to a system, they can constrain the functionality of existing
264 components that they rely on. To address this tension, we have continuously
265 worked to decouple Autopilot into subcomponents with clear inheritance hierar-
266 chies and APIs that can be used quasi-independently.

267 Modularity has 3 primary advantages:

- 268 1. **Modularity makes code more flexible** by reducing the constraints imposed by
269 unstructured code interdependencies
- 270 2. **Modularity makes code more intelligible** by logically distributing tasks to
271 discrete classes
- 272 3. **Modularity reduces effort-duplication** by allowing multiple, similar classes
273 to be created with inheritance rather than copying and pasting.



274 There is no such thing as “incompatible hardware” with Autopilot because the classes
275 that control hardware are independent from the code that provides other core func-
276 tionality. In systems without modular design, hardware implementation is spread
277 across the codebase. For example to add a new type of hardware output to a Bpod
278 system, one would need to write new firmware for it in C (eg. the **valve driver mod-**
279 **ule**), **modify Bpod's existing firmware**, hunt through the code to modify how **states**
280 **are added** and **state machines are assembled**, add its controls explicitly **to the GUI**,
281 and so on.

282 Tasks specify what type of hardware is needed to run them, but are agnostic about
 283 the way the hardware is implemented, making their descriptions more portable. Tasks
 284 that have the same structure but differ in hardware (eg. a freely moving two-alternative
 285 forced choice task in which a mouse visits several IR sensors, or a head-fixed two-
 286 alternative forced choice task in which a mouse runs on a wheel to indicate its choice)
 287 can be implemented by a trivial subclass that modifies the hardware description
 288 rather than completely rewriting the task.

289 *Plugins & Code Transparency*

290 We call Autopilot a software framework because in addition to providing classes
 291 and methods to run experiments out of the box, it also provides explicit structure
 292 that scaffolds any additional code that is needed by the user. Our goal is to clearly
 293 articulate in the documentation how modules should interact so that anyone can
 294 write code that works on any apparatus.

295 As groupware intended to be used differently by lab members with different respon-
 296 sibilities, Autopilot is designed for users with a range of programming expertise,
 297 from those who only want to interact with a GUI, to those who wish to fundamen-
 298 tally rewrite core operations for their particular experiment. As such, it is extensivly
 299 documented: this paper provides a high-level introduction to its design and struc-
 300 ture, its user guide describes how to use the program and provides examples, and
 301 its API-level documentation describes in granular detail how the program actually
 302 works¹. Nothing is “off-limits” to the user—there isn’t any hidden, undocumented
 303 hardware code behind the curtain². We want users to be able to understand how
 304 and why everything works the way it does so that Autopilot can be adapted and
 305 expanded to any use-case.

306 A broader goal of Autopilot is to build a library of flexible task prototypes that can
 307 be tweaked and adapted, hopefully reducing the number of times the wheel is rein-
 308 vented. We have attempted to nudge users to write reusable tasks by designing Au-
 309 topilot such that rather than writing tasks as local unstructured scripts, they use its
 310 plugin system that scaffolds development by extending any of its basic types. Plug-
 311 ins are registered using a form in the Autopilot Wiki which makes them available to
 312 anyone while also embedding them in a semantically annotated information system
 313 that allows giving explicit credit to contributors, programmatically linking to any
 314 derivative publications that use the plugin, and further documentation of any tasks,
 315 hardware, or other extensions included within the plugin. Inheriting from parent
 316 classes give plugins structure and a set of basic features³ while also being maximally
 317 permissive — anything can be overridden and modified.

318 *Message Handling*

319 Modular software needs a well-defined protocol to communicate between modules,
 320 and Autopilot’s is heavily influenced by the concurrency philosophy⁴ of ZeroMQ[30]. All communication between computers and modules happens with ZeroMQ mes-
 321 sages, and handling those messages is the main way that Autopilot handles events. A
 322 key design principle is that Autopilot components should not “share state”—they
 323 can communicate, but they are not *dependent* on one another. While this may seem
 324 like a trivial detail, having networking and message-handling at its core has three ad-
 325 vantages that make Autopilot a fundamental departure from previous behavioral

¹ The user guide and API documentation are available at docs.auto-pi-lot.com

² For readability of the docs, we omit generating HTML documentation for some private methods and functions, but they are documented in the source and their function is made clear from their context and the documentation of public methods.

³ Like inheriting from the `GPIO` class gives GPIO plugins a systematic means of interacting with the underlying pigpio daemon.

⁴ “ZeroMQ [...] has a subversive effect on how you develop network-capable applications. [...] message processing rapidly becomes the central loop, and your application soon breaks down into a set of message processing tasks.”

“If there’s one lesson we’ve learned from 30+ years of concurrent programming, it is: *just don’t share state.*”

327 software.

328 First, new software modules can be added to any system by simply dropping in a
 329 standalone networking object. There is no need to dramatically reorganize existing
 330 code to make room for new functionality. Instead new modules can receive, pro-
 331 cess, and send information by just connecting to another module in the swarm. For
 332 example, each `plot` opens a network connection to stream incoming task data inde-
 333 pendently from the stream that is saving the data.

334 Second, Autopilot can be made to interact with other software libraries that use
 335 ZeroMQ. For example, The OpenEphys GUI for electrophysiology [can send and](#)
 336 [receive ZMQ messages](#) to execute actions such as starting or stopping recordings.
 337 Interaction with other software is also useful in the case that some expensive com-
 338 putation needs to happen mid-task. For example, one could send frames captured
 339 from a video camera on a Raspberry Pi to a GPU computing cluster for tracking
 340 the position of the animal. Since ZeroMQ messages are just TCP packets it is also
 341 possible to communicate over the internet for remote control or to communicate
 342 with a data server.

343 Third, making every component network-capable allows tasks to be distributed over
 344 multiple Raspberry Pis. Chaining multiple Pis distributes the computational load,
 345 allowing, for example, one Raspberry Pi to record and process video while another
 346 runs a sound server and delivers rewards. Autopilot expands with the complexity
 347 of your task, simultaneously eliminating limitations on quantity of hardware per-
 348 ipherals while ensuring latency is minimal. More interestingly, distributing tasks
 349 allows the arbitrary construction of what we call “behavioral topologies,” which we
 350 describe in [section 3.7](#).

351 2.2 Reproducibility

352 We take a broad view on reproducibility: including not only the ability to share data
 353 and recreate experiments, but also integrating into a broader ecosystem of tools that
 354 reduces labor duplication and encourages sharing and organizing technical knowl-
 355 edge. For us, reproducibility means building a set of tools that make every experi-
 356 ment and every technique available to anyone, anywhere.

357 Standardized task descriptions

358 The implementation and fine details of a behavioral experiment matter. Seemingly
 359 trivial details like milliseconds of delay between trial phases and microliters of re-
 360 ward volume can be the difference between a successful and unsuccessful task (Figure
 361 [2.1](#)). *Reporting* those details can thus be the difference between a reproducible
 362 and unreproducible result. Researchers also often use “auxiliary” logic in tasks—
 363 such as methods for correcting response bias—that are never completely neutral for
 364 the interpretation of results. These too can be easily omitted due to brevity or mem-
 365 ory in plain-English descriptions of a task, such as those found in Methods sections.
 366 Even if all details of an experiment were faithfully reported, the balkanization of be-
 367 havioral software into systems peculiar to each lab (or even to individuals within
 368 a lab) makes actually performing a replication of a behavior result expensive and
 369 technically challenging. Widespread use of experimental tools that are not explic-
 370 itly designed to preserve every detail of their operation presents a formidable barrier

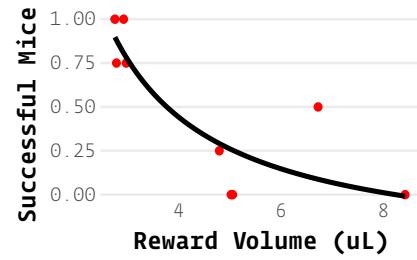


Figure 2.1: “Minor” details have major effects. Proportion of mice (each point, n=4) that were successful learning the first stage of the speech task described in [31] across 10 behavior boxes with variable reward sizes. A 2 μ L difference in reward size had a surprisingly large effect on success rate.

- 371 to rigorous and reproducible science[14].
- 372 Autopilot splits experiments into a) the **code** that runs the experiment, which is
373 intended to be standardized and shared across implementations, and b) the **param-
374 eters** (Figure 2.2) that define your particular experiment and system configuration.
375 For example, two-alternative forced choice tasks have a shared structure regardless
376 of the stimulus modality, but only your task plays pitch-shifted national anthems.
377 This division of labor, combined with Autopilot’s structured plugin system, help
378 avoid the ubiquitous problem of rig-specific code and hard-coded variables making
379 experimental code only useful on the single rig it was designed for — enabling the
380 possibility of a shared library of tasks as described in section 2.1.
- 381 The practice of reporting exactly the parameter description used by the software to
382 run the experiment removes any chance for incompleteness in reporting. Because all
383 task parameters are included in the produced data files, tasks are fully portable and
384 can be reimplemented exactly by anyone that has comparable hardware to yours.

385 *Self-Documenting Data*

- 386 A major goal of the open science movement is to normalize publishing well doc-
387 umented and clearly formatted data alongside every paper. Typically, data are ac-
388 quired and stored in formats that are lab-idiosyncratic or ad-hoc, which, over time,
389 sprout entire software libraries needed just to clean and analyze it. Idiosyncratic
390 data formats hinder collaboration within and between labs as the same cleaning and
391 analysis operations gain multiple, mutually incompatible implementations, dup-
392 licating labor and multiplying opportunities for difficult to diagnose bugs. Over time
393 these data formats and their associated analysis libraries can mutate and become
394 incompatible with prior versions, rendering years of work inaccessible or uninter-
395 pretiable. In one worst-case scenario, the cleaning process unearths some critically
396 missing information about the experiment, requiring awkward caveats in the Meth-
397 ods section or months of extra work redoing it. In another, the missing information
398 or bugs in analysis code are never discovered, polluting scientific literature with in-
399 accuracies.
- 400 The best way to make data publishable is to avoid cleaning data altogether and *design
401 good data hygiene practices into the data acquisition process*. Autopilot automatically
402 stores all the information required to fully reconstruct an experiment, including any
403 changes in task parameters or code version that happen throughout training as the
404 task is refined.

- 405 Autopilot data is stored in **HDF5** files, a hierarchical, high-performance file format.
406 HDF5 files support metadata throughout the file hierarchy, allowing annotations
407 to natively accompany data. Because HDF5 files can store nearly all commonly used
408 data types, data from all collection modalities—trialwise behavioral data, continu-
409 ous electrophysiological data, imaging data, etc.—can be stored together from the
410 time of its acquisition. Data is always stored with the full conditions of its collection,
411 and is ready to analyze and publish immediately (Figure 2.3). No Autopilot-specific
412 scripts are needed to import data into your analysis tool of choice—anything that
413 can read HDF5 files can read Autopilot data⁵.
- 414 As of v0.5.0, we have built a formal data modeling system into Autopilot, allowing
415 for unified declaration of data for experimental subjects, task parameters, and re-
416 sulting data with verifiable typing and human-readable annotations. These abstract

```
{
  "step_name"      : "tone_discrim",
  "task_type"     : "2AFC",
  "bias_mode"     : 0,
  "punish_sound"  : false,
  "stim" : {
    "sounds" : {
      "L": {
        "duration"   : 100,
        "frequency" : 10000,
        "type"       : "tone",
        "amplitude"  : 0.01},
      "R": {"... ":"..."}}
    },
  "reward": {
    "type"      : "volume",
    "volume"   : 20},
  "graduation": {
    "type"      : "accuracy",
    "threshold" : 0.75,
    "window"    : 400},
}
```

Figure 2.2: Task parameters are stored as portable JSON, formatting has been abbreviated for clarity.

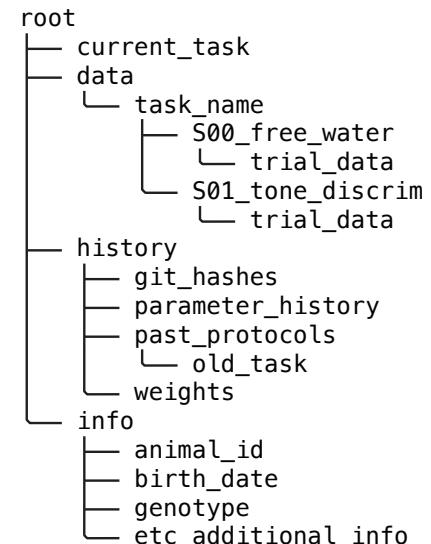


Figure 2.3: Example data structure. All information necessary to reconstruct an experiment is automatically stored in a human-readable HDF5 file.

⁵ Though our `Subject` class provides a simplified interface to access and manipulate Autopilot data

417 data models can be used with multiple storage interfaces, paving the way for export
 418 to, for example, the Neurodata Without Borders standard[32], further enabling Au-
 419 topilot data to be immediately incorporated into existing processing pipelines (see
 420 section 3.2).

421 *Testing & Continuous Integration*

422 Open-source scientific software does away with prior limitations to access and in-
 423 spection imposed by proprietary tools. It also exposes the research process to bugs
 424 in software written by semi-amateurs that can yield errors in the resulting data, anal-
 425 ysis, and interpretation[33, 34, 35, 36]. Autopilot tries to bring best practices in
 426 software development to experimental software, including a set of automated tests
 427 for continuous integration.

428 We are still formalizing our contribution process, and our tests are still far from
 429 achieving full coverage⁶, but we currently require tests and documentation for all
 430 new code added to the library. Writing good tests is hard, and we are in the process
 431 of building a set of hardware simulators and test fixtures to ease contribution.

432 Tests are effectively provable statements about how a program functions (Figure
 433 2.4), which are particularly important for a library that aspires to be baseline lab
 434 infrastructure like Autopilot. Tests make it possible to use and contribute to the li-
 435 brary with confidence: all tests are run on every commit, making it possible to deter-
 436 mine if some new contribution breaks existing code without manually reading and
 437 testing every line. As we work to complete our test coverage, we hope to provide
 438 researchers with a tool that they can trust and elevates the verifiability of scientific
 439 results at large.

440 *Expense*

441 Autopilot is an order of magnitude less expensive than comparable behavioral sys-
 442 tems (Table 2.1). We think the expense of a system is important for two reasons:
 443 scientific equity and statistical power.

444 The distribution of scientific funding is highly skewed, with a large proportion of
 445 research funding concentrated in relatively few labs[37]. Lower research costs ben-
 446 efit all scientists, but lower instrumentation costs directly increase the accessibility
 447 of state-of-the-art experiments to labs with less funding. Since well-funded labs also
 448 tend to be concentrated at a few (well-funded) institutions, lower research costs also
 449 broaden the base of scientists outside traditional research institutions that can stay
 450 at the cutting edge[38, 39, 40].

⁶ Coverage statistics for Autopilot are available on coveralls.io at <https://coveralls.io/github/auto-pi-lot/autopilot>

```
def test_set_gpio():
    """
    The `set` method of a Digital_Out
    object sets the pin state
    """
    pin = Digital_Out(pin=17)

    # Turn GPIO pin on
    pin.set(True)
    assert pin.state == True

    # Turn GPIO pin off
    pin.set(False)
    assert pin.state == False
```

Figure 2.4: A test like `test_set_gpio` is a provable statement about the functionality of a program, in this case that “the `Digital_Out.set()` method sets the state of a GPIO pin.”

451 Neuroscience also stands to benefit from the lessons learned from the replication
 452 crisis in Psychology[42]. In neuroscience, underpowered experiments are the rule,
 453 rather than the exception[43]. Statistical power in neuroscience is arguably even
 454 worse than it appears, because large numbers of observations (eg. neural record-
 455 ings) from a small number of animals are typically pooled, ignoring the nested struc-
 456 ture of observations collected within individual animals. Increasing the number of
 457 cells recorded from a small number of animals dramatically increases the likelihood
 458 of Type I errors (Figure 2.5)—indeed, for values of within-animal correlation typi-
 459 cal of neuroscientific data, high numbers of observations make Type I errors more
 460 likely than not[41]. For this reason, perhaps paradoxically, recent technical advances
 461 in multiphoton imaging and silicon-probe recordings will actually make statistical
 462 rigor in neuroscience *worse* if we don’t use analyses that account for the multilevel
 463 structure of the data and correspondingly record from the increased number of an-
 464 imals that they require.

465 Although the expense of multi-photon imaging and high-density electrophysiology
 466 will always impose an experimental bottleneck, behavioral training time is often the
 467 greater determinant of study sample size. Typical behavioral experiments require
 468 daily training sessions often carried out over weeks and months, while far fewer imag-
 469 ing or electrophysiology sessions are carried out per animal. Training large cohorts
 470 of animals in parallel is thus the necessary basis of a well-powered imaging or elec-
 471 trophysiology experiment.

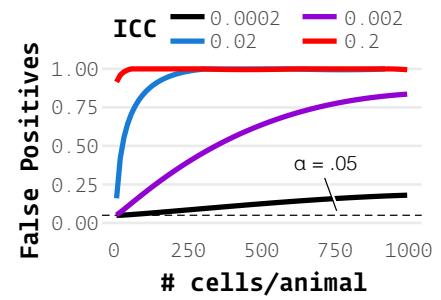


Figure 2.5: When comparing a value across groups, eg. a genetic knockout vs. wildtype, even a modest intra-animal (or, more generally, intra-cluster) correlation (ICC) causes the false positive rate to be far above the nominal $\alpha = 0.05$. Shown are false positive rates for simulated data with various numbers of “cells” recorded for comparisons between two groups of 5 animals each with a real effect size of 0. We note that 741 simultaneously recorded cells were reported in [7] and a mean ICC of 0.19 across 18 neuroscientific datasets was reported in [41]

	Autopilot	pyControl	Bpod
Behavior CPU	\$45	\$270	\$925
Nosepoke (3x)	\$216	\$369	\$810
Total for One	\$261	\$639	\$1735
Five Systems	\$1305	\$3195	\$8675
Host CPU(s)	\$1000	\$1000	\$5000
Total for Five	\$2305	\$4195	\$13625
Total for Ten	\$3610	\$8390	\$27350

Table 2.1: Cost for Basic 2AFC System

“Nosepoke” includes a solenoid valve, IR sensor, water tube, LED, housing, and any necessary driver PCBs. For PyControl and Autopilot, we included the cost of one Lee LHDA0531115H solenoid valve per nosepoke (\$63.35). For PyControl, we estimated a typical USB hub with 5 ports to control 5 pyControl systems from one computer. We note that the Bpod and PyControl systems both include cost of assembly for the control CPUs and nosepokes, but also that Autopilot does not require assembly for its control CPU and its default nosepoke is a snap-together 3D printed part and PCB without surface mounted components that can be assembled by an amateur in roughly half an hour.

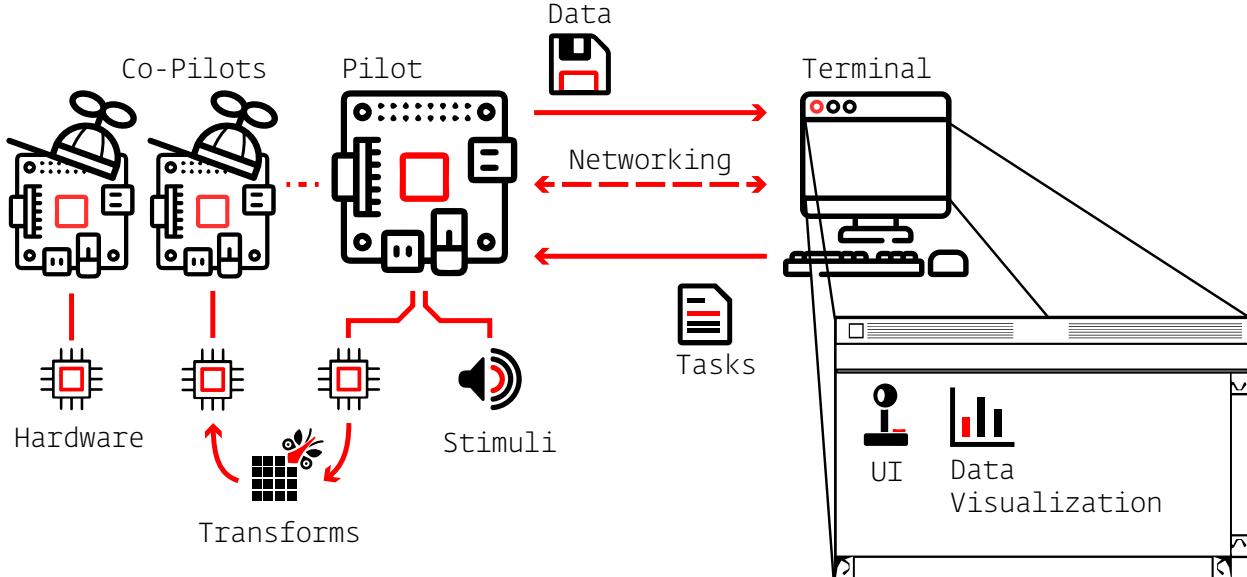
473 *Methods & Results*

Figure 3.1: Overview of major Autopilot components

474 AUTOPILOT CONSISTS OF SOFTWARE AND HARDWARE MODULES
 475 configured to create a behavioral **topology**. Independent **agents** linked by flexible
 476 **networking** objects fill different roles within a topology, such as hosting the **user**
 477 **interface**, controlling **hardware**, **transforming** incoming and outgoing data, or de-
 478 **livering** **stimuli**. This infrastructure is ultimately organized to perform a behavioral
 479 **task**. Autopilot distributes experiments across a network of Raspberry Pis,¹ a type
 480 of inexpensive single-board computer.

¹ Raspberry Pi model 4B, see Table 3.2

481 *3.1 Design*

482 Autopilot has three primary design principles:

- 483 1. **Efficiency** - Autopilot should minimize computational overhead and maximize
 484 use of hardware resources.
- 485 2. **Flexibility** - Autopilot should be transparent in all its operations so that users
 486 can expand it to fit their existing or desired use-cases. Autopilot should provide
 487 clear points of modification and expansion to reduce local duplication of labor
 488 to compensate for its limitations.
- 489 3. **Reproducibility** - Autopilot should maximize system transparency and mini-
 490 mize the potential for the black-box of local reprogramming. Autopilot should
 491 maximize the information it stores about its operation as part of normal data
 492 collection.

493 *Efficiency*

494 Though it is a single board, the Raspberry Pi operates more like a computer than a
 495 microcontroller. It most commonly runs a custom Linux distribution, Raspbian,
 496 allowing Autopilot to use Python across the whole system. Using an interpreted
 497 language like Python running on Linux has inherent performance drawbacks com-
 498 pared to compiled languages running on embedded microprocessors. In practice
 499 these drawbacks are less profound than they appear on paper: Python's overhead is
 500 negligible on modern processors², jitter and performance can be improved by wrap-
 501 ping compiled code, etc. While we view the gain in accessibility and extensibility
 502 of a widely used high-level language like Python as outweighing potential per-
 503 formance gains from using a compiled language, Autopilot is nevertheless designed to
 504 maximize computational efficiency.

505 Most behavioral software is single-threaded (Figure 3.2), meaning the program will
 506 only perform a single operation at a time. If the program is busy or waiting for an
 507 input, other operations are blocked until it is finished.

508 Autopilot distributes computation across multiple processes and threads to take ad-
 509 vantage of the Raspberry Pi's four CPU cores. Most operations in Autopilot are
 510 executed in **threads**. Specifically, Autopilot spawns separate threads to process mes-
 511 sages and events, an architecture described more fully in section 3.8. Threading does
 512 not offer true concurrency³, but does allow Python to distribute computational
 513 time between operations so that, for example, waiting for an event does not block
 514 the rest of the program, and events are not missed because the program is busy (Fig-
 515 ure 3.3).

516 Critical operations that are computationally intensive or cannot be interrupted are
 517 given their own dedicated **processes**. Linux allows individual cores of a processor to
 518 be reserved for single processes, so individual Raspberry Pis are capable of running
 519 four truly parallel processing streams. For example, all Raspberry Pis in an Autopilot
 520 swarm create a messaging client to handle communication between devices which
 521 runs on its own processor core so no messages are missed. Similarly, if an experiment
 522 requires sound delivery, a realtime **sound engine** in a separate process (Figure 3.4)
 523 also runs on its own core.

524 Since even moderately complex experiments can consume more resources than are
 525 available on a single processor, the topmost layer of concurrency in Autopilot is
 526 to use additional **computers**. Autopilot uses the Raspberry Pi as a low-cost hard-
 527 ware controller, but only its GPIO control system is unique to them: the rest of
 528 the code can be used on any type of computer, so computationally expensive or
 529 GPU-intensive operations can be offloaded to any number of high performance ma-
 530 chines. Computers divide labor *autonomously* (see 2.1 and 3.7), so for example one
 531 computer running a task can send and receive messages from another running the
 532 GUI and plots, but does not *depend* on that input as it would in a system that cou-
 533 ples a microcontroller with a managing computer. The ability to coordinate multi-
 534 ple, autonomous computers with heterogeneous responsibilities and capabilities in
 a shared task is Autopilot's definitive design decision.

² and improvements to CPython in Python 3.11 and onwards will bring overhead close to zero[44]

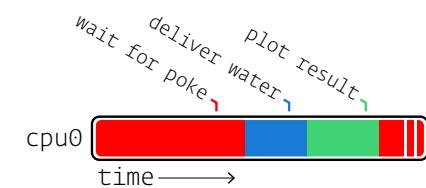


Figure 3.2: A single-threaded program executes all operations sequentially, using a single process and cpu core.

³ See David Beazley's 'Understanding the Global Interpreter Lock' and associated visualizations.

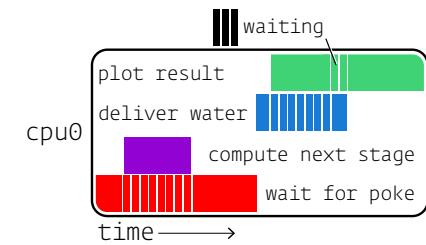


Figure 3.3: A multi-threaded program divides computation time of a single process and cpu core across multiple operations so that, for example, waiting for input doesn't block other operations.

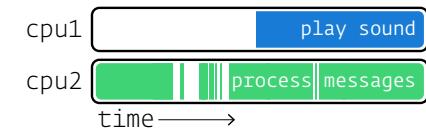


Figure 3.4: A multi-process program is truly concurrent, allowing multiple cpu cores to operate in parallel.

536 *3.2 Data*

537 As of v0.5.0, Autopilot uses [pydantic](#) to create explicitly typed and schematized
538 data models. Submodules include `data` abstract modeling tools that define base
539 model types like `Tables`, `Groups`, and sets of `Attributes`. These base modeling
540 classes are then built into a few core data models like `subject` `Biography` informa-
541 tion, `Protocol` declaration, and the `Subject` data model itself that combines them.
542 Modeling classes then have multiple `interfaces` that can be used to create equiva-
543 lent objects in other formats, like `pytables` for `hdf5` storage, `pandas` dataframes for
544 analysis, or exported to Neurodata Without Borders.

545 For example, consider a simplified version of the Biography model:

Listing 1: data - Biography

```

1 from autopilot.data.modeling import Data, Attributes, Field
2 from typing import Optional, Union
3 from datetime import datetime, timedelta
4
5 class Enclosure(Data):
6     """Where does the subject work?"""
7     box: Optional[Union[str, int]] = Field(
8         default=None,
9         description="The box this Subject is run in")
10    room: Optional[Union[str, int]] = Field(
11        default=None,
12        description="The room number that the animal is run in")
13
14 class Biography(Attributes):
15     """Biography of an Experimental Subject"""
16     id: str = Field(...,
17                     description="The identifying name of this subject.")
18     dob: datetime = Field(...,
19                     description="The Subject's date of birth")
20     enclosure: Optional[Enclosure] = None
21
22 @property
23 def age(self) -> timedelta:
24     """Difference between now and :attr:`.dob`"""
25     return datetime.now() - self.dob

```

Data models use builtin Python type hints. Type hints are colon delimited annotations like `x:int` that indicate the type (integer, string, etc.) of the variable. Though typically Python does not, Pydantic both validates that a type matches its hint and coerces it to the correct type if possible.

The `Union` type means that a field can be one of several possible types, in this case the box can be identified with either a string or integer.

Optional fields can have default fields, either a single value like `None` or a function that computes a default value like the current date.

Descriptions are stored in the data model schema to make shared data self-documenting, and also used by GUI widgets for tooltips that clarify what fields mean.

The class that our model inherits from indicates how Autopilot should treat it in a given storage interface. For HDF5 files, subclasses of the `Attributes` class are stored as node metadata, while subclasses of `Table` make tables.

Autopilot's format interfaces define mappings from nonstandard types to types supported by the format. For HDF5 files, `datetime` objects are converted to ISO 8601 formatted strings

Data models can be recursive, or use other models as types for their own fields. In this case the subject's Enclosure can be optionally specified in its Biography.

Properties are model fields that are automatically computed based on the values of other fields. The age of the animal can be accessed like a normal instance attribute that returns a `timedelta` object.

546 A new subject could then be created with a biography like this, storing it in the HDF5 file and made accessible through the Subject interface:

Listing 2: data - New Subject

```

1 from autopilot.data import Subject
2 from autopilot.data.models import Biography, Enclosure
3 bio = Biography(
4     id="my_subject",
5     dob="2022-01-01T00:00:00",
6     enclosure=Enclosure(box=100, room="Building 200"))
7
8 sub = Subject.new(bio)
9 assert sub.info == bio

```

The `Subject` class is the primary means by which Autopilot stores, organizes, and interacts with data.

Several basic models are built into Autopilot, and in future versions it will be possible to extend and replace these models with plugins, making storage formats fully customizable while still being explicit and understandable.

If we don't give the type specified in the model, it will try and coerce it to the correct type and raise an error if it can't.

An `assert` declares that some logical statement is `True` and raises an exception if it isn't. Autopilot's unit tests ensure that subject data can be stored and retrieved without losing information or changing types.

548 The models are declared using a combination of python type hints and `Field` objects
 549 that provide defaults and descriptions. Because these models can be recursive,
 550 as in the case of using the `Enclosure` model as a type within the `Biography` model,
 551 we can build expressive, flexible, but still strict representations of complex data.

552 Out of the box, pydantic models can create explicit and interoperable `schemas` in
 553 `JSON Schema` and `OpenAPI` formats, and Autopilot extends them with additional
 554 interfaces and representations. Autopilot can create a GUI form for filling in fields
 555 for models, for example, to create a new `Subject` or declare parameters for a task
 556 ([Figure 3.5](#)). Attribute models that consist of scalar key-value pairs can be reliably
 557 stored and retrieved from metadata attribute sets in HDF5 groups, but Autopilot
 558 knows that `Table` models should be created as HDF5 tables as they will have mul-
 559 tiple values for each field. An additional `Trial_Data` class that inherits from `Table`
 560 can be exported to NWB trial data, and the `Subject.get_trial_data` method uses
 561 the model to load trial data and convert it to a correctly typed pandas[[45](#)] `DataFrame`.

562 Though the data modeling system is new in v0.5.0⁴, we have laid the groundwork
 563 for Autopilot’s plugin system to allow researchers to declare custom schema for all
 564 data produced by Autopilot, and to preserve both interoperability and reproducibil-
 565 ity by combining them with datasets potentially produced by multiple incompatible
 566 tools (see [Section 5.4](#)).

567 3.3 Tasks

568 Behavioral experiments in Autopilot are centered around `tasks`. Tasks are Python
 569 classes that describe the parameters, coordinate the hardware, and perform the logic
 570 of the experiment. Tasks may consist of one or multiple `stages` like a stimulus pre-
 571 sentation or response event, completion of which constitutes a `trial` ([Figure 3.6](#)).
 572 Stages are analogous to states in the finite state machine formalism.

573 Multiple tasks are combined to make `protocols`, in which animals move between
 574 tasks according to “graduation” criteria like accuracy or number of trials. Training
 575 an animal to perform a task typically requires some period of shaping where they
 576 are familiarized to the apparatus and the structure of the task. For example, to teach
 577 animals about the availability of water from “nosepoke” sensors, we typically begin
 578 with a “free water” task that simply gives them water for poking their nose in them.
 579 Having a structured protocol system prevents shaping from relying on intuition or
 580 ad hoc criteria.

581 Task Components

582 The following is a basic two-alternative choice (2AFC) task—a sound is played and
 583 an animal is rewarded for poking its nose in a designated target nosepoke. While
 584 simple, it is included here in full to show how one can program a task, including an
 585 explicit data and plotting structure, in roughly 60 lines of generously spaced Python.

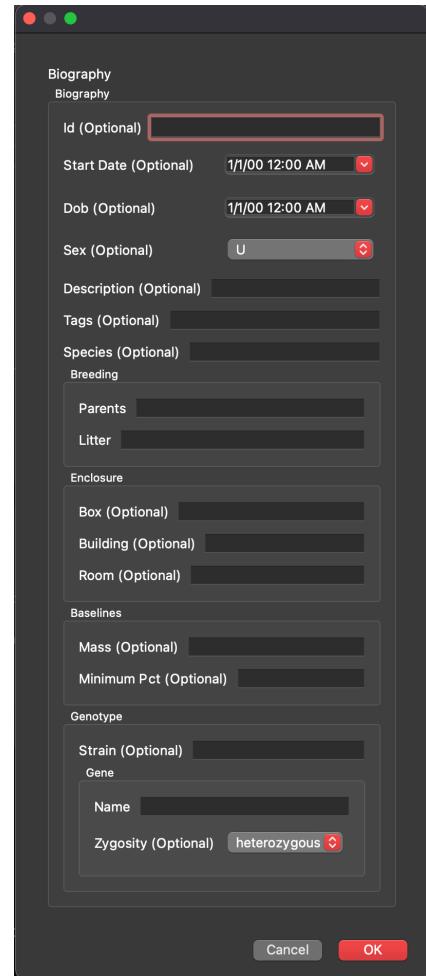


Figure 3.5: An Autopilot Data model can automatically generate a GUI form to fill in its properties, in this example to define a new experimental Subject’s biography.

⁴ Released as an alpha version at the time of writing

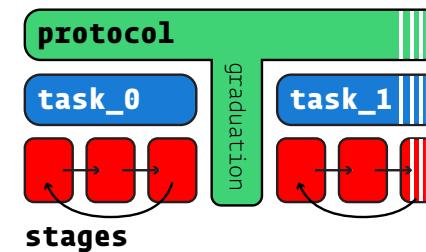


Figure 3.6: Protocols consist of one or multiple tasks, tasks consist of one or multiple stages. Completion of all of a task’s stages constitutes a trial, and meeting some graduation criterion like accuracy progresses a subject between tasks.

586 Every task begins by describing four elements:

587 1) the task's parameters, 2) the data that will be collected, 3) how to plot the data,
588 and 4) the hardware that is needed to run the task.

Listing 3: task - parameters

```

1  class Nafc(Task):
2      class Params(Task_Parms):←
3          stim: Sounds = Field(...,
4              description = "Sound Stimuli")
5          reward: units.mL = Field(...,
6              description= "Reward Volume (mL)")
7
8
9  class TrialData(Trial_Data): ←
10     target: Side = Field(...,
11         description="Side (L, R) of the correct response")
12     correct: bool = Field(...,
13         description="Response matched target")
14
15     PLOT = {} ←
16     PLOT['data'] = {'target' : 'point',
17                      'correct' : 'rollmean'},
18     # n trials to roll window over
19     PLOT['params'] = {'roll_window' : 50}
20
21     HARDWARE = { ←
22         'POKES':{
23             'L': 'Digital_In',
24             'R': 'Digital_In'
25         },
26         'PORTS':{
27             'C': 'Solenoid', ←
28         }
29     }

```

1) A **Task_Parms** model defines what parameters are needed to run the task.

We use **Field** objects as in listing 1, and can also use some special types like **Sounds** to declare complex parameters

units work like numbers but avoid ambiguity, so eg. the **Solenoid** class below knows this is a volume, rather than a duration

2) A **Trial_Data** model defines what data will be returned from the task.

3) A **PLOT** dictionary maps the data output to graphical elements in the GUI. (In future versions this will be incorporated into the **Fields** of **TrialData**)

4) A **HARDWARE** dictionary that describes what hardware will be needed to run the task.

The specific implementation of the hardware (eg. where it is connected, how to interact with it) is independent of the task. The task just knows about a PORT named 'C' that is a **Solenoid**.

589

590 Created tasks receive some common methods, like input/trigger handling and net-
591 working, from an inherited metaclass. Python inheritance can also be used to make
592 small alterations to existing tasks⁵ rather than rewriting the whole thing. The GUI
593 will use the **Params** model and the **PLOT** dictionary to generate forms for parameter-
594 izing the task within a protocol and display the data as it is collected. The **Subject**
595 class will use the **TrialData** model to create HDF5 tables to store the data, and the
596 **Task** metaclass will instantiate the described **HARDWARE** objects from their system-
597 specific configuration in the **prefs.json** file so they are available in the rest of the
598 class like **self.hardware['POKES']['L'].state**

⁵ An example of subclassing a generic 'Task' class is included in Autopilot's [user guide](#)

599 *Stage Methods*

600 The logic of tasks is described in one or a series of methods (stages). The order of
 601 stages can be cyclical, as in this example, or can have arbitrary logic governing the
 602 transition between stages.

Listing 4: task - methods

```

30 def __init__(self, params:'Nafc.Params'):
31     self.stim_mgr = Stim_Manager(params.stim)
32     self.reward   = Reward_Manager(params.reward)
33
34     stage_list  = [self.discrim, self.reinforcement]
35     self.stages = itertools.cycle(stage_list)
36
37     self.init_hardware()
38     next(self.stages)() ←
39
40 def discrim(self):
41     target, wrong, stim = self.stim_mgr.next() ←
42     self.target = target
43
44     self.triggers[target] = [
45         self.hardware['PORTS']['C'].open, ←
46         lambda: next(self.stages)()] ←
47     self.triggers[wrong] = lambda: next(self.stages)()
48
49     self.node.send('DATA', {'target':target}) ←
50
51     stim.play()
52
53 def reinforcement(self, response): ←
54     if response == self.target:
55         self.node.send('DATA', {'correct':True})
56     else:
57         self.node.send('DATA', {'correct':False})
58
59     next(self.stages)() ←

```

- In Python, `def` defines new methods. The `__init__` method is called when a new object is initialized
- Managers control stimulus and reward delivery, so users can, for example, continually synthesize new stimuli or implement adaptive rewards
- Stages are combined into an object that (in this case) continually cycles through them when its `next()` method is called.
- This starts the task by retrieving the first stage and then calling it.
- The stimulus manager returns which port will be the target and the sound to be played.
- A sequence of triggers is set: if the target port is poked, a reward will be delivered and the next stage will be called. A `lambda` function indicates not to call the method now, but only when triggered.
- The task has a networking object that asynchronously streams data back to the user-facing terminal
- In this example, the response port is passed from the trigger handling function. If it matches the stored target variable, the animal answered correctly.
- Finally, the task is repeated by calling the next stage.

603 Autopilot is not prescriptive about how tasks are written. The same task could have
 604 two separate methods for correct and incorrect answers rather than a single reinforcement method, or only a single stage that blocks the program while it waits for
 605 a response.

607 Publishing data from this task requires no additional effort: a hash that uniquely
 608 identifies the code version (as well as any local changes) is automatically stored at
 609 the time of collection, as is a JSON-serialized version of the parameter model (Figure 3.7). If this task was incorporated into the central task library, anyone using Autopilot would be able to exactly replicate the experiment from the published data.

```
{
  "step_name": "Simple 2AFC",
  "stim": {
    "sounds": {
      "L": {
        "type": "tone",
        "frequency": 4000
      },
      "R": {
        "type": "tone",
        "frequency": 8000
      }
    },
    "reward": 10
}
```

Figure 3.7: Simplified example of parameters for the above task

612 *The limitations of finite state machines*

613 The 2AFC task described above could be easily implemented in a finite-state ma-
 614 chine. However, the difficulty of programming a finite-state machine is subject to
 615 combinatoric explosion with more complex tasks. Specifically, finite-state machines
 616 can't handle any task that requires any notion of "state history."

617 As an example, consider a maze-based task. In this task, the animal has to learn a
 618 particular route through a maze—it is not enough to reach the endpoint, but the
 619 animal has to follow a specific path to reach it (Figure 3.8). The arena is equipped
 620 with an actimeter that detects when the animal enters each area.

621 In Autopilot, we would define a hardware object that logs positions from the actime-
 622 ter with a `store_position()` method. If the animal has entered the target position
 623 ("i" in this example), a `task_trigger()` that advances the task stage is called. The
 624 following code is incomplete, but illustrates the principle.

Listing 5: `maze - hardware`

```

1  class Actimeter(Hardware):
2      def __init__(self):
3          # ... some code to access the hardware ...
4          self.positions = []
5          self.target_position = "i"
6
7      def store_position(self, position):
8          self.positions.append(position)
9
10     if position == self.target_position:
11         self.finished_cb(self.positions) ←
12         self.positions = []

```

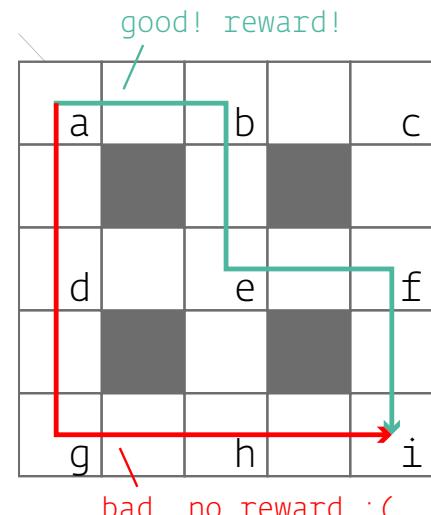


Figure 3.8: The subject must reach point i but only via the correct (green) path.

See line 18 below

625 The task follows, with parameters and network methods for sending data omitted
 626 for clarity.

Listing 6: `maze - task`

```

13  class Maze(Task):
14      def __init__(self):
15          self.target_path = ['a', 'b', 'e', 'f', 'i']
16
17          self.actimeter = Actimeter()
18          self.actimeter.finished_cb = self.finished ←
19
20      def finished(self, positions):
21          if positions == self.target_path: ←
22              self.reward()

```

The actimeter is given a reference to the Maze task's `finished()` method, which it calls when the target position is reached

The sequence of `positions` is compared to the `target_path` with `=`. If they match, the subject is rewarded!

628 How would such a task be programmed in a finite-state machine formalism? Since
 629 the path matters, each “state” needs to consist of the current position and all the
 630 positions before it. But, since the animal can double back and have arbitrarily many
 631 state transitions before reaching the target corner, this task is impossible to represent
 632 with a finite-state machine, as a full representation would necessitate infinitely many
 633 states (this is one example of the *pumping lemma*, see [46]).

634 Even if we dramatically simplify the task by 1) assuming the animal never turns back
 635 and visits a space twice, and 2) only considering paths that are less than or equal to
 636 the length of the correct path, the finite state machine would be as complex as figure
 637 3.9.

638 While finite-state machines are relatively easy to implement and work well for simple
 639 tasks, they quickly become an impediment to even moderately complex tasks. Even
 640 for 2AFC tasks, many desirable features are difficult to implement with a finite state
 641 machine, such as: (1) graduation to a more difficult task depending on performance
 642 history, (2) adjusting reward volume based on learning rate, (3) selecting or synthe-
 643 sizing upcoming stimuli based on patterns of errors[47], etc.

644 Some of these problems are avoidable by using extended versions of finite state ma-
 645 chines that allow for extra-state logic, but require additional complexity in the code
 646 running the state machines to accomodate, and with enough exceptions the clean
 647 systematicity that is the primary benefit of finite state machines is lost. Autopilot
 648 attempts to avoid these problems by providing *tools* to program tasks and describe
 649 them without *requiring a specified format*, balancing the increased complexity by
 650 scaffolding the broader ecosystem of the experiment like its output data, hardware
 651 control, etc. When possible, we have tried to avoid forcing people to change the way
 652 they think about their work to fit our “little universe”⁶ and instead try to provide a
 653 set of tools that let researchers decide how they want to use them.

⁶ We take inspiration from Aaron Swartz’ de-
 scription of another engineering project, the
 Semantic Web, that became too precious about its
 formalisms:

“Instead of the “let’s just build something
 that works” attitude that made the Web (and the
 Internet) such a roaring success [...] they formed
 committees to form working groups to write
 drafts of ontologies that carefully listed (in 100-
 page Word documents) all possible things in the
 universe and the various properties they could
 have, and they spent hours in Talmudic debates
 over whether a washing machine was a kitchen
 appliance or a household cleaning device. [...] And
 instead of spending time building things, they’ve
 convinced people interested in these ideas that the
 first thing we need to do is write *standards*. (To
 engineers, this is absurd from the start—standards
 are things you write *after* you’ve got something
 working, not before!)”[48]

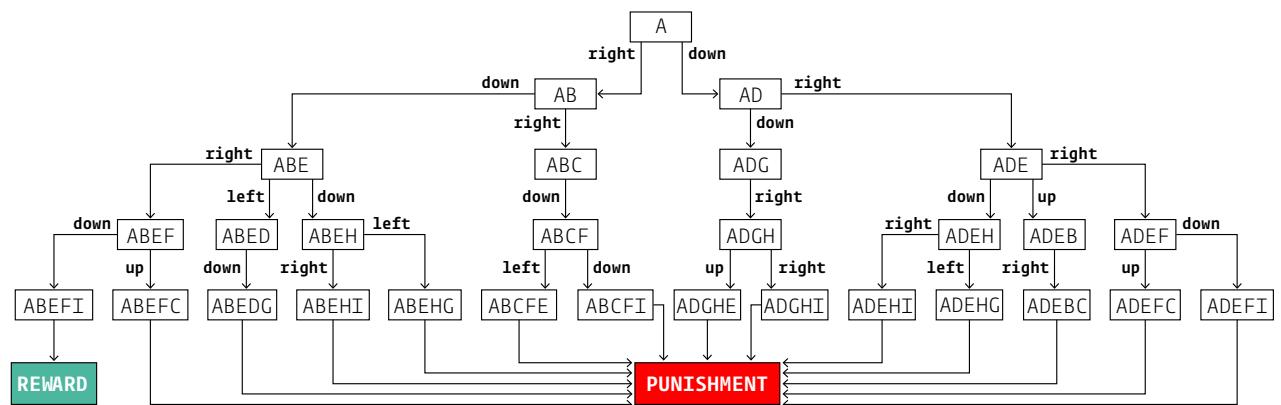


Figure 3.9: State transition tree for a simplified maze task.

654 3.4 Hardware

655 The Raspberry Pi can interface with nearly all common hardware, and has an [extensive collection](#) of [guides](#), [tutorials](#), and an active [forum](#) to support users implementing new hardware. There is also an enormous amount of existing hardware for the Raspberry Pi, including [sound cards](#), [motor controllers](#), [sensor arrays](#), [ADC/DACs](#), and [touchscreen displays](#), largely eliminating the need for a separate ecosystem of purpose-built hardware (Table 3.1).

661 Autopilot controls hardware with an extensible inheritance hierarchy of Python
662 classes intended to be built into a library of hardware controllers analogously to
663 tasks. Autopilot uses [pigpio](#) to interact with its GPIO pins, giving Autopilot 5 μ s
664 measurement precision and enabling protocols that require high precision (such as
665 Serial, PWM, and I2C) for nearly all of the pins. Currently, Autopilot also has a
666 family of objects to control cameras (both the [Raspberry Pi Camera](#) and [high-speed](#)
667 [GENICAM-compliant](#) cameras), i2c-based [motion](#) and [heat](#) sensors, and [USB mice](#).
668 In the [future](#) we intend to improve performance further by replacing time-critical
669 hardware operations with low-level interfaces written in Rust.

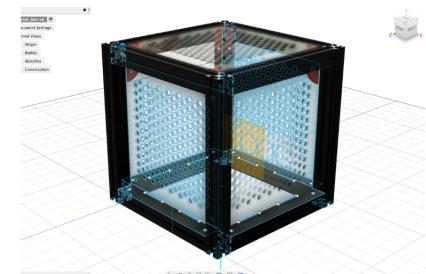
670 To organize and make available the vast amount of contextual knowledge needed to
671 build and use experimental hardware, we have made a densely linked and publicly
672 editable [semantic wiki](#). The Autopilot wiki contains, among others, reference in-
673 formation for [off-the-shelf parts](#), schematics for [2D](#) and [3D](#)-printable components,
674 and [guides](#) for building experimental apparatuses and custom parts. The wiki com-
675 bines unrestricted freeform editing with structured, computer-readable [semantic](#)
676 [properties](#), and we have defined a collection of [schemas](#) for commonly documented
677 items coupled with [submission forms](#) for ease of use. For example, the wiki page
678 for the [Lee Company solenoid](#) we use has fields from a generic [Part](#) schema like a
679 datasheet, price, and voltage, but also that it's a 3-way, normally-closed [solenoid](#).

680 The wiki's blend of structure and freedom breaks apart typically monolithic hard-
681 ware documentation into a collaborative, multimodal technical knowledge graph.
682 Autopilot can access the wiki through its [API](#), and we intend to tighten their integra-
683 tion over time, including automatic configurations for common parts, usage and
684 longevity benchmarks, detecting mutually incompatible parts, and automatically re-
685 solving any additional plugins or dependencies needed to use a part.

Table 3.1: Cost of common peripherals. The native hardware of the Raspberry Pi, low-level hardware control of Autopilot, and availability of inexpensive off-the-shelf components compatible with the raspi make most custom-built peripherals unnecessary.

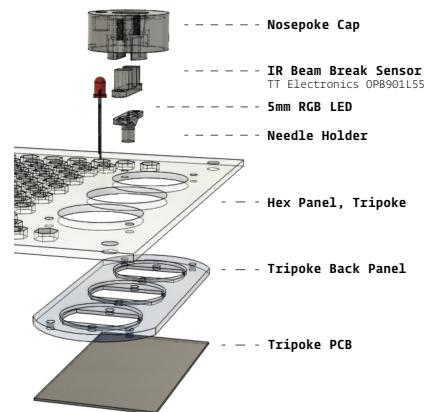
Device	Raspi	Bpod
HiFiBerry DAC2 Pro	\$45	\$445
ADC	\$30	\$495
I2C	\$0	\$225
Ethernet	\$0	\$285
Rotary Encoder	\$0	\$145

Autopilot Behavior Box



A modular box for training mouse-sized animals

Modality	Enclosures
Build Guide Type	Construction Build Guide
Creator	Jonny Saunders
Version	2
Submitted Date	2021-06-10



Autopilot Tripoke

Figure 3.10: Two examples of parts with assembly guides available on the autopilot wiki: A modular behavior box with magnetic snap-in panels (top), and a three-nosepoke panel (bottom).

Table 3.2: Specifications of reviewed behavior hardware. BPod's state machine uses the Teensy 3.6 microcontroller, and PyControl uses the Micropython Pyboard.

	Raspberry Pi 4B	Teensy 3.6	pyboard
CPU Clock	1.5GHz	180MHz	168MHz
CPU Cores	4	1	1
Architecture	ARMv8-A, 64-bit	ARMv7 32-bit	ARMv7 32-bit
RAM Size	2, 4, or 8GB	256KB	192KB
Storage	MicroSD (any size)	1024KB	1024KB
GPU	Broadcom VideoCore VI	—	—
GPIO Pins	40	58	29
USB Ports	2x USB 2.0, 2x USB 3.0	2x USB 2.0	1x USB 2.0
Ethernet	1Gbps	100Mbps	—
WiFi	2.4/5 GHz b/g/n/ac	—	—
Camera	15-pin Serial Interface	—	—
Bluetooth	✓	—	—

686 3.5 Transforms

687 In v0.3.0, we introduced the `transform` module, a collection of tools for transforming
 688 data off a sensor. The raw data off a sensor is often not in itself useful for performing an
 689 experiment: we want to compare it to some threshold, extract positions of objects
 690 in a video feed, and so on. Transforms are like building blocks, each performing
 691 some simple operation with a standard object structure, and then composed into a
 692 pipeline (Figure 3.11). Pipelines are portable, and can be created on the fly from a
 693 JSON representation of their arguments, so it's easy to offload expensive operations
 694 to a more capable machine for distributed realtime experimental control (See [29]).

695 In addition to computing derived values, we use transforms in a few ways, including

- 696 • **Bridging Hardware** — Different hardware devices use different data types,
 697 units, and scales, so transforms can `rescale` and convert values to make them compatible.
- 699 • **Integrating External Tools** — The number of exciting analytical tools for real-
 700 time experiments keep growing, but in practice they can be hard to use together.
 701 The transform module gives a scaffolding for writing wrappers around other
 702 tools and exposing them to each other in a shared framework, as we did with
 703 DeepLabCut-Live[29], making closed-loop pose tracking available to the rest of
 704 Autopilot's ecosystem. We don't need to rally thousands of independent develop-
 705 ers to agree to write their tools in a shared library, instead transforms make
 706 wrapping them easy.
- 707 • **Extending Objects** — Transforms can be used to augment existing objects and
 708 create new ones. For example, a `motion sensor` uses the `spheroid` transform to
 709 calibrate its accelerometer, and the `gammatone filter`⁷ extends the `Noise` sound
 710 to make a gammatone `filtered noise` sound.

711 Like Tasks and Hardware, the transform module provides a scaffolding for writing
 712 reference implementations of algorithms commonly needed for realtime behavioral
 713 experiments. For example, neuroscientists often want to quickly measure a research
 714 subject's velocity or orientation, which is possible with inexpensive inertial motion
 715 sensors (IMUs), but since anything worth measuring will be swinging the sensor
 716 around with wild abandon the readings first need to be rotated back to a geocentric
 717 coordinate frame. Since the readings from an accelerometer are noisy, we found a
 718 few whitepapers describing using a Kalman filter for fusing the accelerometer and
 719 gyroscope data for a more accurate orientation estimate ([49, 50]), but couldn't find
 720 an implementation. We `wrote one` and integrated it into the IMU class (Figure 3.12).
 721 Since it's an independent transform, it's available to anyone even if they use nothing
 722 else from Autopilot.

723 Transforms were made to be composed, so we broke it into independent sub-operations:
 724 A `Kalman` filter, `rotation`, and a `spheroid correction` to calibrate accelerometers. Then
 725 we combined it with the DLC-Live transform for a fast but accurate motion esti-
 726 mate from position, velocity, and acceleration measurements from three indepen-
 727 dent sensors. Since each step of the transformation is exposed in a clean API, it was
 728 straightforward to `extend the Kalman filter` to accomodate the the wildly different
 729 sampling rates of the camera and IMU. It's still got its quirks, but that's the pur-
 730 pose of plugins — to make the code `available and documented` without formally
 731 integrating it in the library.



New in
v0.3.0:
Transform

```
— Transform - DLC Live! —
from autopilot import transform as t
# track points on a human body
dlc = t.image.DLC(
    model_zoo="full_human")
# select one of the knees
dlc += t.selection.DLCSlice(
    select="knee2")
# Test if it's in an ROI
dlc += t.logical.Condition(
    minimum=(0,0),
    maximum=(128,128))

# Process frames with the pipeline
# set pin High if knee2 in ROI
while True:
    pin.set(dlc.process(
        cam.q.get()))
```

Figure 3.11: Transforms can be chained together (here with the in-place addition operator `+=`) to make pipelines that encapsulate the logical relationship between some input and a desired output. Here `pin` is a `Digital_Out` object, and `cam` is a `PiCamera` with queue enabled.

⁷a thin wrapper around `scipy's signal.gammatone`

```
— Geocentric Velocity —
# rotate input in x and y
# by some pitch and roll
reorient = t.geometry.Rotate(
    dims='xy')
# select the z axis
reorient += t.selection.Slice(
    select=2)
# remove gravity
reorient += t.math.Add(
    -9.8)

# using I2C_9DOF ...
angle = imu.rotation
z_accel = reorient.process(
    imu._acceleration, angle)
```

Figure 3.12: Using the `IMU_Orientation` transform built into the IMU's `rotation` property, a processing chain to reorient the accelerometer reading and subtract gravity for geocentric z-axis acceleration.

732 **3.6 Stimuli**

733 A hardware object would control a speaker, whereas stimulus objects are the individual sounds that the speaker would play. Like tasks and hardware, Autopilot makes
 734 stimulus generation portable between users, and is released with a family of common sounds like tones, noises, and sounds from files. The logic of sound presentation is contained in an inherited metaclass, so to program a new stimulus a user
 735 only needs to describe how to generate it from its parameters (Figure 3.13). Sound stimuli are better developed than visual stimuli as of v0.5.0, but we present a proof-of-concept visual experiment (Section 4.5) using `psychopy`[21].

741 Autopilot controls the realtime audio server `jack` from an independent Python process
 742 that dumps samples directly into jack’s buffer (Figure 3.14), giving it a trigger-to-playback latency very near the theoretical minimum (Section 4.2). Sounds can
 743 be pre-buffered in memory or synthesized on demand to play continuous sounds.
 744 Because the realtime server is independent from the logic of sound synthesis and
 745 storage, stimuli can be controlled independently from different threads without interrupting audio or dropping frames.

748 We use the [Hifiberry Amp 2](#), a combined sound card and amplifier, which is capable of 192kHz/24Bit audio playback. Autopilot and Jack can output to any sound hardware, however, including the builtin audio of the Raspberry Pi if fidelity isn’t important. There are no external video cards for the Raspberry Pi 4b⁸, but its embedded video card is capable of presenting video and visual stimuli (Section 4.5) especially if the other computationally demanding parts of the task are distributed to other Raspberry Pis (Section 3.7). If greater video performance is needed, Autopilot is capable of running on typical desktops as well as other single-board computers with GPUs (as we did with the Nvidia Jetson in [29]).

757 **3.7 Agents**

758 All of Autopilot’s components can be organized into a single system as an “agent,”
 759 the executable that coordinates everyday use. An agent encapsulates:

- 760 • **Runtime Logic** — an initialization routine that starts any needed system processes and any subsequent operations that define the behavior of the agent.
- 762 • **Networking Station** — Agents have networking objects called Stations that are intended to be the “load bearing” networking objects (described more below).
- 765 • **Callbacks** — An action vocabulary that maps different types of messages to methods for handling them. Called `listens` to disambiguate from other types of callbacks.
- 768 • **Dependencies** — Required packages, libraries, and system reconfigurations needed to operate. Python dependencies are currently defined for agents as groups of optional packages⁹, and system configuration is done with `scripts` which shorthand common operations like `compiling OpenCV` with optimizations for the raspi or enabling a soundcard.

773 Together, these define an agent’s *role* in the swarm.

```
____ An Autopilot Tone _____
my_tone = sounds.Tone(
    frequency = 500,
    duration = 200)
my_tone.play()
```

Figure 3.13: Autopilot stimuli are parametrically defined and inherit all the playback logic that makes them easy to integrate in tasks

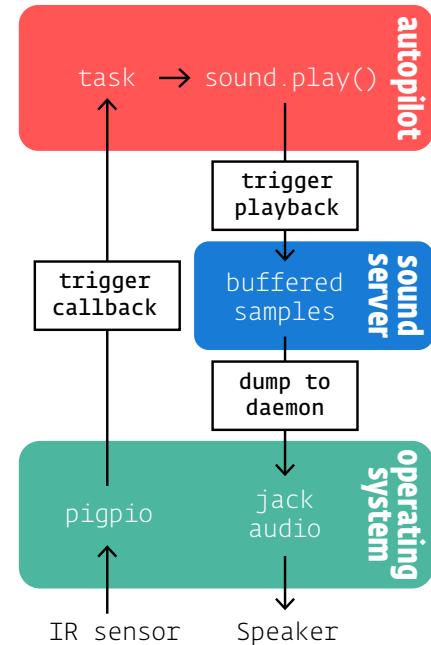


Figure 3.14: Our sound server keeps audio samples buffered until a `.play()` method is called, and then dumps them directly into the jack audio daemon.

⁸ though the Raspberry Pi compute module has a PCI lane that supports GPUs.

⁹ As of v0.5.0, Autopilot is packaged with Poetry, so they are `[tool.poetry.extras]` entries within the `pyproject.toml` file, installed with pip like `pip install auto-pi-lot[pilot]` or poetry like `poetry install -E pilot`

⁷⁷⁴ There are currently two agents in Autopilot:

- ⁷⁷⁵ • **Terminal** - The user-facing control agent.
- ⁷⁷⁶ • **Pilot** - A Raspberry Pi that runs tasks, coordinates hardware, and optionally co-
ordinates a set of child Pis.

⁷⁷⁸ **Terminal** agents serve as a root node (see Section 3.8) in an Autopilot swarm. The
779 terminal is the only agent with a **GUI**, which is used to control its connected pilots
780 and visualize incoming task data. The terminal also manages data and keeps a reg-
781 istry of all active experimental subjects. The terminal is intended to make the day-to-
782 day use of an Autopilot swarm manageable, even for those without programming
783 experience. The terminal GUI is described further in Section 3.9.

⁷⁸⁴ **Pilot** agents are the workhorses of Autopilot—the agents that run the experiments.
785 Pilots are intended to operate as always-on, continuously running system services.
786 Pilots make a network connection to a terminal and wait for further instructions.
787 They maintain the system-level software used for interfacing with the hardware con-
788 nected to the Raspberry Pi, receive and execute tasks, and continually return data
789 to the terminal for storage.

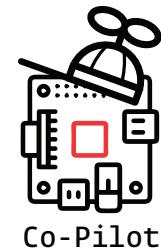
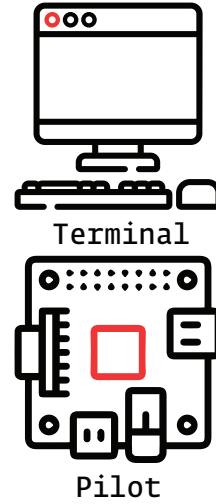
⁷⁹⁰ Each agent runs autonomously, and so a Pilot can run a task without a Terminal
791 and store data locally, a Terminal can be used without Pilots to define protocols and
792 manage subjects, and so on. This decoupling lets each agent have more freedom in
793 its behavior at the expense of the complexity of configuring and maintaining them
794 (see Sec. 5.10 and 5.13). All interaction is based on the “listen” callbacks known by
795 the agents, so to start a task a Terminal will send a Pilot a “START” message contain-
796 ing information about a Task class that it is to run along with its parameterization.
797 The Pilot then attempts to run the task, sends a message to the Terminal alerting it
798 to a “STATE” change, and begins streaming data back to it in messages with a “DATA”
799 key.

⁸⁰⁰ Each pilot is capable of mutually coordinating with one or many **Copilots**¹⁰. We are
801 still experimenting with, and thus openminded to the best way to structure multi-
802 pilot tasks. Like many things in Autopilot, there is no one right way to do it, and the
803 strategy depends on the particular constraints of the task. We include a few examples
804 in the network latency and go/no-go tasks in the **plugin** that accompanies this paper,
805 and expand on this a bit further in a few parts of section 5, as it is a major point of
806 active development.

⁸⁰⁷ *Behavioral topologies*

⁸⁰⁸ We think one of the most transformative features of Autopilot’s distributed struc-
809 ture is the control that users have over what we call “behavioral topology.” The logic
810 of hardware and task operation within an agent, the distribution of labor between
811 agents performing a task, and the pattern of connectivity and command within a
812 swarm of agents constitute a topology.

⁸¹³ Below we illustrate this idea with a few examples:



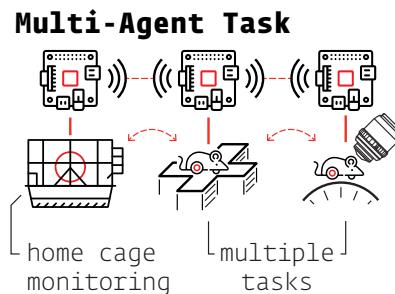
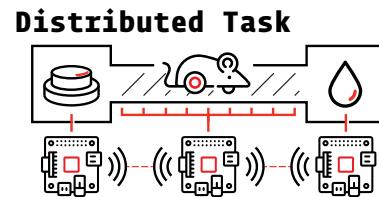
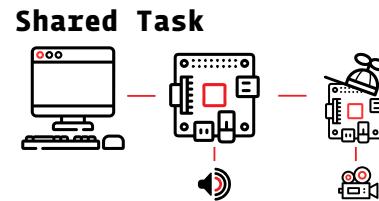
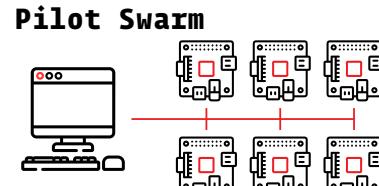
Co-Pilot

¹⁰ A previous version of this paper described a third, subordinate “Child” agent that performed auxiliary operations in a task. We now view such a hierarchy as unnecessary, and that distribution of labor within a task is better served by a fluid combination of multiple Pilots than thinking of them as qualitatively different agents. We now refer one among multiple agents performing a task together as a “copilot.”

- 814 • **Pilot Swarm** - The first and most obvious topological departure from traditional
 815 behavioral instrumentation is the use of a single computer to independently coordinate tasks in parallel. Our primary installation of Autopilot is a
 816 cluster of 10 behavior boxes that can independently run tasks dispatched from
 817 a central terminal which manages data and visualization. This topology highlights
 818 the expandability of an Autopilot system: adding new pilots is inexpensive,
 819 and the single central terminal makes controlling experiments and managing data
 820 simple.
- 822 • **Shared Task** - Tasks can be shared across a set of copilots to handle tasks with
 823 computationally intensive operations. For example, in an open-field navigation
 824 task, one pilot can deliver position-dependent sounds while another records and
 825 analyzes video of the arena to track the animal's position. The terminal only
 826 needs to be configured to connect to the parent pilot, but the other copilot is
 827 free to send data to the Terminal marked for storage in the subject's file as well.
- 828 • **Distributed Task** - Many pilots with overlapping responsibilities can cooperate
 829 to perform distributed tasks. We anticipate this will be useful when the exper-
 830 imental arenas can't be fully contained (such as natural environments), or when
 831 experiments require simultaneous input and output from multiple subjects. Dis-
 832 tributed tasks can take advantage of the Pi's wireless communication, enabling,
 833 for example, experiments that require many networked cameras to observe an
 834 area, or experiments that use the Pis themselves as an interface in a multisubject
 835 augmented reality experiment.
- 836 • **Multi-Agent Task** - Neuroscientific research often consists of multiple mutually
 837 interdependent experiments, each with radically different instrumentation.
 838 Autopilot provides a framework to unify these experiments by allowing users
 839 to rewrite core functionality of the program while maintaining integration be-
 840 tween its components. For example, a neuroethologist could build a new "Ob-
 841 server" agent that continually monitors an animal's natural behavior in its home
 842 cage to calibrate a parameter in a task run by a pilot. If they wanted to manip-
 843 ulate the behavior, they could build a "Compute" agent that processes Calcium
 844 imaging data taken while the animal performs the task to generate and admin-
 845 ister patterns of optogenetic stimulation. Accordingly, passively observed data
 846 can be combined with multiple experimental datasets from across the subject's
 847 lifespan. We think that unifying diverse experimental data streams with interop-
 848 erable frameworks is the best way to perform experiments that measure natural
 849 behavior in the fullness of its complexity in order to understand the naturally
 850 behaving brain[26].

851 3.8 Networking

852 Agents use two types of object to communicate with one another: core **station**
 853 objects and peripheral **node** objects (Figure 3.15). Each agent creates one station
 854 in a separate process that handles all communication *between* agents. Stations are
 855 capable of forwarding data and maintaining agent state so the agent process is not
 856 unnecessarily interrupted. Nodes are created by individual modules run within an
 857 agent—eg. tasks, plots, hardware—that allow them to send and receive messages
 858 within an agent, or make connections directly to other nodes on other agents after
 859 the station discovers their network addresses. Messages are TCP packets¹¹, so there



¹¹ Autopilot uses ZeroMQ[30] and tornado to send and process messages

860 is no distinction between sending messages within a computer, a local network, or
 861 over the internet¹².

862 Both types of networking objects are tailored to their hosts by a set of callback functions — **listens** — that define how to handle each type of message. Messages have
 863 a uniform key-value structure, where the key indicates the listen used to process
 864 the message and the value is the message payload. This system makes adding new
 865 network-enabled components trivial:

866

Listing 7: A new networked LED

```

1 class LED_RGB(Hardware):
2     def __init__(self):
3         # call self.color for a 'COLOR' message
4         self.listens = {'COLOR': self.color}
5         self.node = networking.Node(
6             id      = 'BEST_LED',
7             listens = self.listens)
8
9     def color(msg):
10        self.set_color(msg.value)
11
12    # elsewhere in the code, we change the color to red!
13    node.send(to='BEST_LED', key='COLOR', value=[255, 0, 0])
  
```

¹² Though automatically configuring the use of faster protocols like IPC for communication within an agent or different backends like redis or gstreamer for data streams that would benefit from them is part of our **development goals**

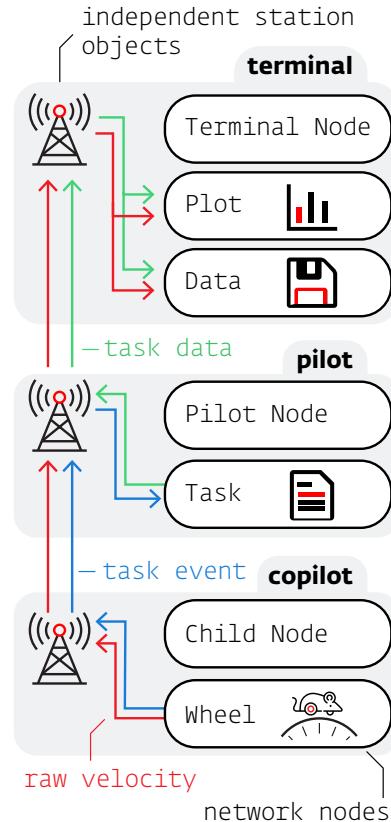


Figure 3.15: Autopilot segregates data streams efficiently—eg. raw velocity (red) can be plotted and saved by the terminal while only the task-relevant events (blue) are sent to the primary pilot. The pilot then sends trial-summarized data to the terminal (green).

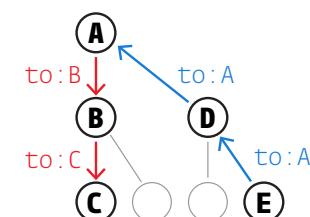


Figure 3.16: Treelike network structure—downstream messages are addressed by successive nodes, but upstream messages can always be pushed until the target is found.

¹³ converted to binary suitable for sending between computers

¹⁴ For example, to send a message from E to C in the diagram above:

```
node.send(to=["A", "B", "C"])
```

888

3.9 GUI & Plots

889 The terminal's GUI controls day-to-day system operation¹⁵. It is intended to be a
 890 nontechnical frontend that can be used by those without programming experience.

891 For each pilot, the terminal creates a control panel that manages subjects, task op-
 892 eration, and plots incoming data. Subjects can be managed through the GUI, in-
 893 cluding creation, protocol assignment, and metadata editing. Protocols can also be
 894 created from within the GUI. The GUI also has a set of basic maintenance and in-
 895 formational routines in its menus, like calibrating water ports or viewing a history
 896 of subject weights.

897 The simple callback design and network infrastructure makes adding new GUI func-
 898 tionality straightforward, and in the future we intend to extend the plugin system
 899 such that plugins can provide additional menu actions, plots, and utilities.

900

Plotting

901 Realtime data visualization is critical for monitoring training progress and ensuring
 902 that the task is working correctly, but each task has different requirements for visu-
 903 alization. A task that has a subject continuously running on a ball might require
 904 a continuous readout of running velocity, whereas a trial-based task only needs to
 905 show correct/incorrect responses as they happen. Autopilot approaches this prob-
 906 lem by assigning the data returned by the task to graphical primitives like points,
 907 lines, or shaded areas as specified in a task's PLOT dictionary (taking inspiration from
 908 Wilkinson's grammar of graphics[51]).

909 The GUI is now some of the oldest code in the library, and we are in the process of
 910 decoupling some of its functionality from its visual representation and moving to
 911 a model where it is a thinner wrapper around the **data modeling tools**. Following
 912 the lead of formal models with strict typing will, for example, make plotting more
 913 fluid where the researcher can map incoming data to the set of graphical elements
 914 that are appropriate for its type. We discuss this further in section 5.8

¹⁵ Autopilot uses **PySide**, a wrapper around **Qt**, to build its GUI.

Trial Plot

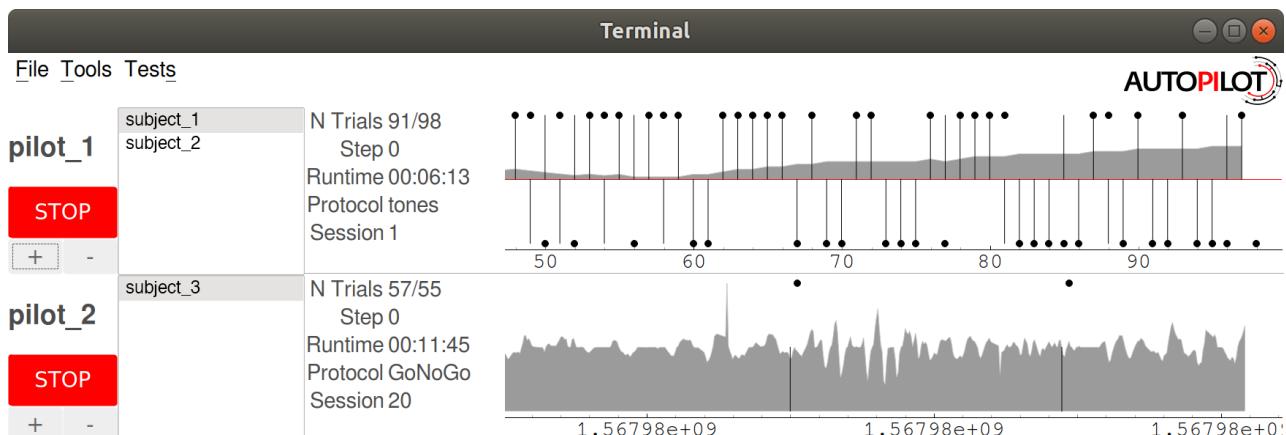
```
{"data": {
    "target" : "point",
    "response" : "segment",
    "correct" : "rollmean"
},
"roll_window" : 50}
```

Continuous Plot

```
{"data": {
    "target" : "point",
    "response" : "segment",
    "velocity" : "shaded"
},
"continuous": true}
```

Figure 3.17: PLOT parameters for Figure 3.18. In both, “target” and “response” data are mapped to “point” and “segment” graphical primitives, but timestamps rather than trial numbers are used for the x-axis in the “continuous” plot (Figure 3.18, bottom). Additional parameters can be specified, eg. the trial plot (Figure 3.18, top) computes rolling accuracy over the past 50 trials

Figure 3.18: Screenshot from a terminal GUI running two different tasks with different plots concurrently. pilot_1 runs 2 subjects: (subject_1 and subject_2), while pilot_2 runs subject_3. See Figure 3.17 for plot description



915

4

916

Tests

917 WE HAVE BEEN TESTING AND REFINING AUTOPILOT since we built our
 918 swarm of 10 training boxes in the spring of 2019. In that time 178 mice¹ have per-
 919 formed over 6 million trials on a range of tasks. While Autopilot is still relatively
 920 new, it is by no means untested.

921 In this section we will present a set of basic performance benchmarks while also
 922 showing several of the different ways that Autopilot can be used. The code for all
 923 of the following tests is available as a [plugin](#) that is further documented on the [wiki](#),
 924 and runs on a prerelease of v0.5.0. Materials tables (Table 4.1) for each test link more
 925 specifically to the test code and provide additional hardware and version documen-
 926 tation, where appropriate.

927

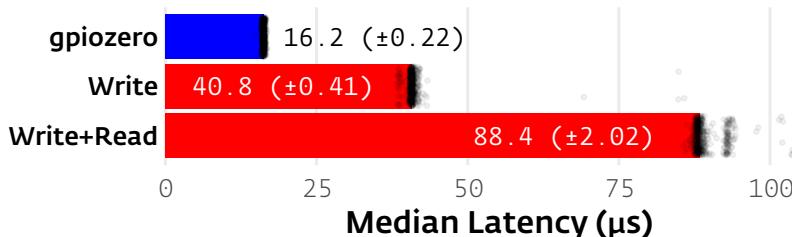
4.1 GPIO Latency

928 Neurons compute at millisecond timescales, so any task that links neural compu-
 929 tation to behavior needs to have near-millisecond latency. We start by characteriz-
 930 ing Autopilot’s GPIO control latency in “script mode” — using the GPIO control
 931 classes on their own, without using any of the rest of Autopilot’s modules.

932

Output Latency

933 We first tested the software measured latency between when a command to write
 934 a value to a GPIO pin is issued and when it completes (Figure 4.1, Table 4.2). By
 935 default, the pigpio interface we use to control GPIO pins issues a command and then
 936 confirms the request was successful by querying the pigpio daemon for the status
 937 of the pin (Write+Read). We [extended](#) pigpio to just issue the command without
 938 confirmation to estimate the true time between when the command is issued and
 939 when the voltage of the pin changes (Write). Each of these operations takes roughly
 940 $40\mu s$ with minimal jitter (Median \pm IQR — Write Only: $40.8\mu s \pm 0.41$, Write and
 941 Read: $88.4\mu s \pm 2.02$, n=100,000 each).



942 Pigpio is useful as a general purpose controller because of its ability to run scripts
 943 within its daemon, use hardware PWM via direct memory access, and consistently
 944 poll for pin state, but takes a latency penalty because the python interface communi-
 945 cates with it through a local TCP socket. To demonstrate the flexibility of Autopilot

¹ All procedures were performed in accordance with National Institutes of Health guidelines, as approved by the University of Oregon Institutional Animal Care and Use Committee.

Table 4.1: General Materials

Hardware	
Raspi	Raspberry Pi 4b
Oscilloscope	Rigol DS1054Z
Software	
Autopilot	v0.5.0a
Plugin	Autopilot_Paper
Python	3.9.12
RaspiOS	Bullseye 22-04-04 (lite)
Analysis	
R	4.2.0
ggplot2 ^[52]	3.3.5
dplyr ^[53]	1.0.9
purrr ^[54]	0.3.4
pandas ^[45]	1.4.2
numpy ^[24, 55]	1.21.6

Table 4.2: GPIO Latency Materials. (Parameters in {} are input in separate runs)

Code	test_gpio.py
replicate	python test_gpio.py -w {0,1,2} -n 100000
gpiozero	1.6.2
pigpio	3c23715

Figure 4.1: Software latency from GPIO write to completion of command. Values are presented as medians \pm IQR with n=100,000 tests for each. A random subsample of 500 (for tractability of plotting) of each type of test are presented (black points) after filtering to the bottom 99th percentile to exclude extreme outliers. Commands sent using pigpio (red) took roughly $40\mu s$ each (write and read are effectively two separate commands, Write Only: $40.8\mu s \pm 0.41$, Write and Read: $88.4\mu s \pm 2.02$). The prototype gpiozero wrapper (blue) using RPi.GPIO as its backend was faster, taking $16.2\mu s \pm 0.22$ to complete.

in incorporating additional software libraries, we wrote a [thin wrapper](#) around `gpi`
`ozero`, which can use `RPi.GPIO` to directly write to the GPIO registers. For simple
output, this wrapper proved to be faster ($16.2 \mu\text{s} \pm 0.22$, n=100,000), and with
63 lines of code is now available in the plugin accompanying this paper to be used,
repurposed, and extended.

951 Roundtrip Input/Output Latency

952 Output commands usually aren't issued in isolation, but as a response to some ex-
953 ternal or task-driven trigger. We measured the roundtrip latency from a 5V square
954 pulse from an external function generator to when an output pin was flipped from
955 low to high on an oscilloscope (Table 4.3).

956 Typically that is as much methodological detail as you would expect in a scientific
957 paper, but actually making those measurements via oscilloscope requires knowing
958 how to set up such a test as well as how to extract the measured data afterwards —
959 which is not altogether trivially available technical knowledge. As an example of how
960 integrating semantically linked documentation with experimental tools enables a
961 fundamentally deeper kind of reproducibility and methodological transparency, we
962 instead documented these operations, including a code sample and a guide to un-
963 locking additional features on our oscilloscope on the [autopilot wiki](#). The code to
964 extract traces from the oscilloscope is also included in this paper's [plugin](#), which
965 links to the oscilloscope page with a [[Controls Hardware :: Rigol DS1054Z]] tag,
966 so it is possible to bidirectionally find code examples from the oscilloscope page as
967 well as find further documentation about the hardware used in this paper from the
968 plugin page. The same can be true for any hardware used by any plugin in any paper
969 using Autopilot.

970 We used the `assign_cb` method of the [Digital_In](#) class to test the typical roundtrip
971 latency that Autopilot objects can deliver. This gave us a median $474 \mu\text{s}$ (IQR:
972 $52.5 \mu\text{s}$) latency (Red in Figure 4.3). GPIO callbacks are flexible, and can use arbi-
973 trary python functions, but if all that's needed is to trigger one pin off of another
974 with some simple logic like a parametric digital waveform or static "on" time, piggpio
975 also allows us to directly program pin to pin logic as a "pigs" script (literally Figure
976 4.2) that runs within the piggpio daemon. The pigs script gave us roughly three or-
977 ders of magnitude lower latency (Median \pm IQR: $370\text{ns} \pm 140$, blue in 4.3).

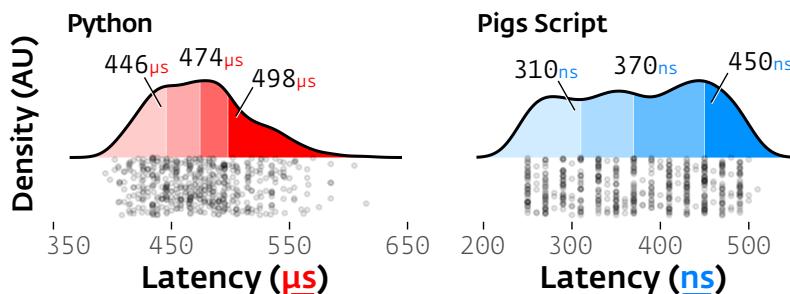


Table 4.3: Roundtrip Latency Materials

Function Generator	Koolertron CJDS98
Code	test_gpio.py
Replicate	python test_gpio.py -w {3,4}

```
____ Pigs Trigger Script ____
" ".join([
    "tag 999",
    # read input pin
    f'r {pin_in.pin_bcm}',
    # if off, goto 998
    f'jz 998',
    # else, turn on
    f'w {pin_out.pin_bcm} 1',
    # then goto 999
    "jp 999",
    "tag 998",
    # turn off
    f'w {pin_out.pin_bcm} 0',
    # jump to beginning
    f"jp 999"
])
```

Figure 4.2: The pigs script used to trigger one pin (pin_out), from another (pin_in). At the expense of a little bit of complexity having to write a script in its scripting language, we are able to reduce latency by three orders of magnitude.

Figure 4.3: Roundtrip latency from external trigger to digital output using two methods: Typical Autopilot callback function given to a `Digital_In` object that turns a `Digital_Out` pin on for 1ms when an input pin changes state (Red, Left, Median \pm IQR: $474 \mu\text{s} \pm 52.5$). Pigs script that runs entirely within the piggpio daemon (Blue, Right, $380\text{ns} \pm 140$). For each, black points represent individual measurements (n=525), annotations are quartiles.

978 4.2 Sound Latency

979 We measured end to end, hardware input to sound output latency by measuring the
 980 delay between an external digital input and the onset of a 10kHz pure tone (Table
 981 4.4). Sound playback was again triggered by the `Digital_In` class's callback method,
 982 and sound samples were buffered in a `deque` held in a separate process by the `jack`
 983 audio client between each trial. A `Digital_Out` pin was wired to the `Digital_In`
 984 pin in order to deliver the trigger pulse (but the `Digital_Out` pin was uninvolved in
 985 the software trigger for sound output).

986 Autopilot's `jack` audio backend was configured with a 192kHz sampling rate with a
 987 buffer with two periods of 32 samples each for theoretical minimum latency of
 988 0.33ms². We observed a median 1.35ms (± 0.72 IQR) latency across 521 samples —
 989 roughly 4x the theoretical minimum (Figure 4.4). This suggests that Autopilot elim-
 990 inates most perceptible end-to-end latency, which is necessary for tasks that require
 991 realtime feedback. One clear future direction is to write the sound processing loop
 992 in a compiled language exposed with a foreign function interface (FFI) to decrease
 993 both latency and jitter.

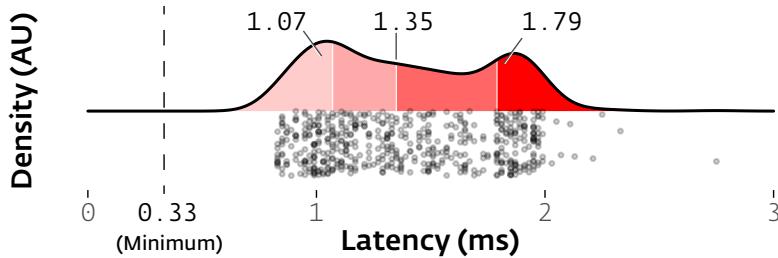


Table 4.4: Sound Latency Materials

Sound Card	Hifiberry Amp2
<code>jack</code>	1.9.22
<code>Code</code>	<code>test_sound.py</code>
<code>Replicate</code>	<code>python test_sound.py</code>

² A previous version of this paper included benchmarking and comparison to Bpod and pyControl's sound onset latency, but since then both packages have changed substantially, including Bpod creating a new `hifi sound module` based off HiFiBerry hardware very similar to the card used here, making those benchmarks obsolete. In this version we have omitted comparative benchmarks in favor of allowing the maintainers of those packages to publish their own benchmarks.

Figure 4.4: Autopilot has a median 1.35ms ($\pm .72$ IQR) latency between an external trigger and sound onset. Individual trials (dots, n=521) are shown beneath a density plot (red area under curve) colored by quartile (shades, numbers above are median, first, and third quartile). This latency is roughly 4x the theoretical minimum (0.33ms, dashed line).

994 4.3 Network Latency

995 To support data-intensive tasks like those that require online processing of video
 996 or electrophysiological data, the networking modules at the core of Autopilot need
 997 high bandwidth and low latency.

998 To test the latency of Autopilot's networking modules, we switch from "script mode"
 999 to "Task mode" (Table 4.5). Tasks are useful for encapsulating multistage routines
 1000 across multiple devices that would be hard to coordinate with scripts alone. Our
 1001 `Network_Latency` task consists of one "leader" pilot sending timestamped messages
 1002 to a "follower" pilot which returns the timestamp marking when it received the mes-
 1003 sage. The two pis communicate via two directly connected `Net_Nodes` (rather than
 1004 routing each message through agent-level `Station` objects) after the leader pi ini-
 1005 tiates the follower with a multihop "START" message routed through a Terminal
 1006 agent containing the task and networking parameters. We measured latency using
 1007 software timestamps while synchronizing the clocks of the two pis with Chrony, an
 1008 NTP daemon previously measured to synchronize Raspberry Pis within dozens of
 1009 microseconds[56]³, with the leader pi hosting an NTP server and the follower pi
 1010 synchronizing its clock solely from the leader. We documented this on the wiki too,
 1011 since synchronization is a universal problem in multi-computer experiments.

1012 Point to point latency was 0.975ms (median, ± 0.1 IQR, n=10,000, Figure 4.5)
 1013 with some clear bimodality where a subset of messages (2,300 of 10,000) took longer

Table 4.5: Network Test Materials

Router	TP-Link AC1750
<code>Chrony</code>	4.0
<code>zmq</code>	22.3.0
<code>Code</code>	<code>Network_Latency</code>
<code>Replicate</code>	Assign task to subject from Terminal, start Task.

³ Our sync is likely to be near to or better than that reported in [56]: in addition to a quiet network, we configured chrony to poll more frequently and tolerate a smaller error than default

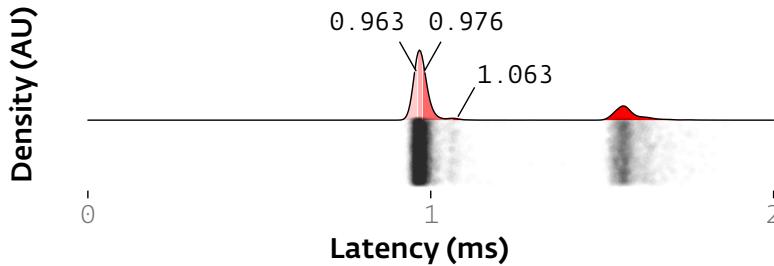


Figure 4.5: Network latency from when a message is sent from one pilot to when it is received by another. Messages took 0.975ms to send and receive (median, ± 0.1 IQR, $n=10,000$, overlaid numbers and red shading in density plot indicate quartiles). There is a clear bimodality in latencies for individual messages (black dots, jittered in y-axis) with unclear cause.

(median 1.567ms). The source of the bimodality is unclear to us, though it could be due to network congestion or interruption by other processes as the networking modules are not run in their own process like the sound server. This latency includes message serialization and deserialization by the builtin JSON library, which is on the order of roughly $100\mu s$ each for even the very small messages sent in this test. In future versions we will explore other serialization tools like `msgpack` and offer them as alternate serialization backends.

1021 4.4 Network Bandwidth

1022 To test Autopilot’s bandwidth, we demonstrate yet another modality of use, using
 1023 Autopilot’s `Bandwidth_Test` widget, an action available from the Terminal GUI’s
 1024 tests menu that corresponds to a callback “listen” method in the Pilot (Table 4.6).
 1025 This test requests that one or several pilots send messages at a range of selected
 1026 frequencies and payload sizes back to the terminal. The messages pass through four
 1027 networking objects en route: the stations and network nodes running the test for
 1028 both the terminal and pilots (See Figure 3.15).

1029 The needs for streaming experimental data vary depending on what is being streamed.
 1030 Electrophysiological data is an n-electrode length vector sampled at a rate of dozens
 1031 to hundreds of kilohertz, so each individual message isn’t very large but there are a
 1032 lot of them. Video data is a width by height (and for color video, by channel) array
 1033 that can be relatively large⁴, but it is captured at dozens to hundreds of hertz. Different
 1034 data streams also have different degrees of compressibility: noisy, quasirandom
 1035 electrical signals compress relatively poorly, while the typical behavioral neurosci-
 1036 entist’s video of an animal that takes up 1/10th of the frame against a white back-
 1037 ground can have compression ratios in the hundreds.

1038 Autopilot tries to provide flexibility for streaming different data types by offering
 1039 message batching and optional on-the-fly compression with `blosc`. The bounds on
 1040 bandwidth are then the speed at which an array can be compressed and the rate at
 1041 which messages of a given size can be sent.

1042 Autopilot’s networking modules were able to send an “empty” (402 byte) message
 1043 with headers describing the test but no payload at a maximum observed rate of
 1044 1,818Hz⁵. Approximately 15% of the duration is spent in message serialization, as a
 1045 “frozen” preserialized message can be sent at 2,100Hz, though we imagine the need
 1046 to send the same message thousands of times is rare.

1047 We tested four types of messages with nonzero array⁶ payloads: since the entropy of
 1048 an array determines how compressible it is, we sent random and all-zero arrays with

Table 4.6: Bandwidth Test Materials.

Terminal	Macbook Pro 2019, macOS 12.3.1, 2.4 GHz 8-Core Intel Core i9 v0.2.0
blosc2	<code>Bandwidth_Test</code> ,
Code	<code>Pilot.l_bandwidth</code>
Replicate	Terminal > Tools > Test Bandwidth

⁴ $(1920 * 1080 * 3 * 8 \text{ bits}) / 8 = 6$ megabytes per frame of a 1080p color video, which is why video is rarely streamed uncompressed

⁵ maximum average rate of 5000 messages for each of the equivalent empty message tests in the four conditions described below

⁶ In all cases, float64 numpy arrays encoded in base64

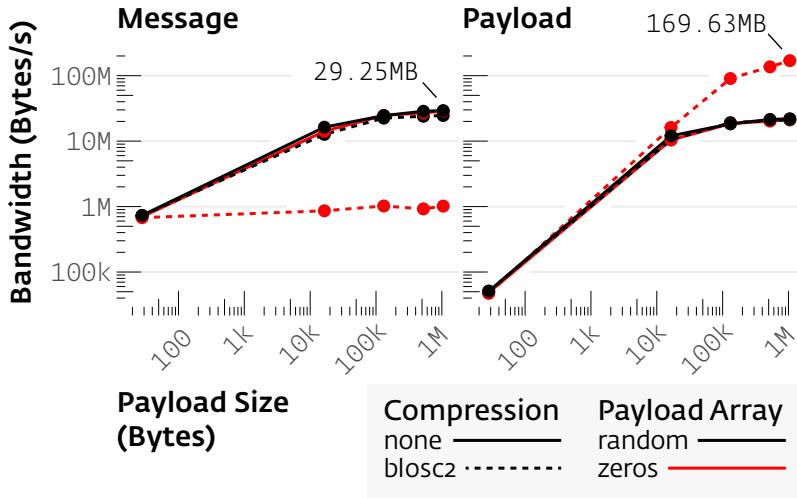


Figure 4.6: Bandwidth measurements between a pilot and terminal for compressed (solid lines) or uncompressed (dotted lines) arrays of random numbers (black) or zeros (red, each point $n=5000$ messages). As message size increased, the bandwidth for the rate of bytes transferred in serialized messages (“message bandwidth,” left) plateaued at 29.25MBytes/s, while the effective bandwidth of arrays before and after compression (“payload bandwidth,” right) reached 169.63MBytes/s. Real data will fall somewhere in this effective bandwidth range, depending on its compressibility.

and without compression. The random and all-zero arrays are the floor and ceiling of compressibility, respectively. Compression gives us two notions of bandwidth: the literal number of bytes that can be passed through a connection, and the effective bandwidth of the size of the arrays that can be transferred with a given compression ratio. We refer to these as “message” and “payload” bandwidth, respectively in Figure 4.6. Message bandwidth reflects the hardware limitations of the Raspberry Pi, but payload bandwidth is the number that matters in practice, as it measures the actual “speed of data” that can be used by the receiver.

As we increased the size of the array payload⁷, the message bandwidth plateaued at a maximum of 29.25MByte per second (Figure 4.6, left). After this plateau, increasing the message size trades off linearly with the rate of messages sent. For all but the compressed array of zeros, the payload bandwidth mirrored the message bandwidth with some trivial overhead from the base64 encoding. The compressed array of zeros, however, had an effective payload bandwidth of 169.6MBytes/s, a compromise between the speed of compression with the smaller message size⁸. The compressed random array had only negligible differences in payload and message bandwidth compared to the uncompressed random array, indicating that the overhead for blosc is trivial.

The ability to batch messages allows researchers to tune the size of an individual message to their particular need for high bandwidth or low latency. Since the compressibility of real data varies across the entire entropic range from randomness to arrays of all zeros, Autopilot doesn’t have a single “bandwidth”, but one that ranges between 30 and 170MByte/s⁹. This bandwidth makes Autopilot capable of streaming raw Calcium imaging¹⁰ and electrophysiological data from modern high-density probes¹¹. Its flexible architecture allows researchers to decide how to build their experiments by distributing different components over different combinations of computers: stream data from a raspberry pi to a more powerful computer for processing, use GPIO rather than network triggers for time-critical operations — meeting the tooling challenge of complex, hardware-intensive, multimodal experiments that define contemporary systems neuroscience.

⁷ $n=5,000$ for each condition at each size

⁸ A message with a 1MByte zero array payload compressed to 6KBytes.

⁹ In this dataset. There is additional payload bandwidth headroom with larger messages, and we include an additional dataset with a 200MByte/s bandwidth in the supplement.

¹⁰ 2-Photon: 5.9MB/s
(12 bits * 512x512 resolution * 15Hz)

¹¹ Neuropixels: 14.4MB/s^[7]
(10 bits * 30kHz * 384 channels)

1079 4.5 Distributed Go/No-go Task

1080 We designed a visual go/no-go task as a proof of concept for distributing task el-
 1081 ements across multiple Pis, and also for the presentation of visual stimuli (Figure
 1082 4.7, Table 4.7). While the rest of the tests presented have been re-run, in the time
 1083 since the initial publication of the preprint we have not done substantial work on
 1084 Autopilot's visual stimulus module, and so this section is presented as previously
 1085 written using the v0.1.0 initial release.

1086 In this task, a head-fixed subject would¹² be running on a wheel in front of a display
 1087 with a lick-detecting water port able to deliver reward. Above the port is an LED.
 1088 Whenever the LED is green, if the subject drops below a threshold velocity for a fixa-
 1089 tion period, a grating stimulus at a random orientation is presented on the monitor.
 1090 After a random delay, there is a chance that the grating changes orientation by a ran-
 1091 dom amount. If the subject licks the port in trials when the orientation is changed,
 1092 or refrains from licking when it is not, the subject is rewarded.

1093 One pilot controlled the operation of the task, including the coordination of a copi-
 1094 lot. The pilot was connected to the LED and solenoid valve for reward delivery, as
 1095 well as a monitor¹³ to display the gratings¹⁴. The copilot continuously streamed
 1096 velocity data (measured with a USB optical mouse against the surface of the wheel)
 1097 back to the terminal for storage (see also Figure 3.15, which depicts the network
 1098 topology for this task). The copilot waited for a message from the pilot to initiate
 1099 measuring velocity, and when a rolling average of recent velocities fell below a given
 1100 threshold the copilot sent a TTL trigger back to the pilot to start displaying the grat-
 1101 ing. This split-pilot topology allows us to poll the subject velocity continuously (at
 1102 125Hz in this example) without competing for resources with psychopy's rendering
 1103 engine.

1104 We measured trigger (TTL pulse from the copilot) to visual stimulus onset latency
 1105 using the measurement cursors of our oscilloscope as before. To detect the onset
 1106 of the visual stimulus, we used a high-speed optical power meter attached to the
 1107 top-left corner of our display monitor. The stimulus was a drifting Gabor grating
 1108 drawn to fill half the horizontal and vertical width of the screen (960 x 540px), with
 1109 a spatial frequency of 4cyc/960px and temporal (drift) frequency of 1Hz.

1110 We observed a bimodal distribution of latencies (Quartiles: 28, 30, 36ms, n=50,
 1111 Figure 4.8), presumably because onsets of visual stimuli are quantized to the refresh
 1112 rate (60Hz, 16.67ms) of the monitor. This range of latencies corresponds to the
 1113 second and third frame after the trigger is sent (2/3 of observations fall in the 2nd
 1114 frame, 1/3 of observations in the 3rd frame). We observed a median framerate of
 1115 36.2 FPS (IQR: 0.7) across 50 trials (8863 frames, Figure 4.9).

1116 We further tested the Pi's framerate by using Psychopy's `timeByFrames` test—a script
 1117 that draws stimuli without any Autopilot components running—to see if the fram-
 1118 erate limits were imposed by the hardware of the Raspberry Pi or overhead from
 1119 Autopilot (Table 4.8). We tested a series of Gabor filters and `random dot` stimuli
 1120 (dots travel in random directions with equal velocity, default parameters) at differ-
 1121 ent screen resolutions and stimulus complexities. The Raspberry Pi was capable
 1122 of moderately high framerates (>60 FPS) for smaller, lower resolution stimuli, but
 1123 struggled (<30 FPS) for full HD, fullscreen stimuli.

Table 4.7: Go/No-go Materials

Hardware	
Beam Break	TT Electronics OPB901L55
Monitor	Acer S230HL
Lick Port	Autopilot Tripoke v1
Light Sensor	Thorlabs PM100D
Software	
psychopy	v3.1.5
glfw	v1.8.3

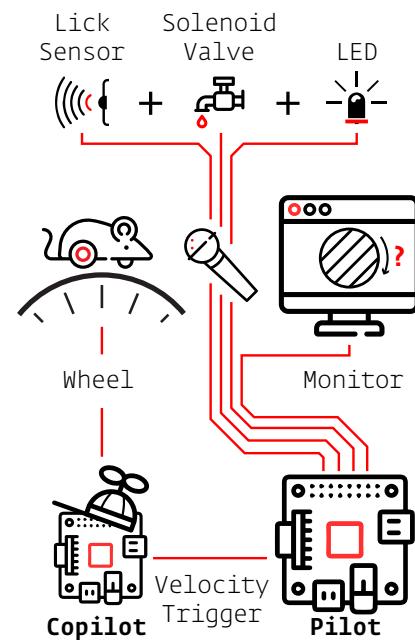


Figure 4.7: Hardware distribution for the distributed go/no-go task. Red lines indicate physical connections between hardware components. The lick sensor, solenoid valve, and LED are physically bundled into one component represented as the mouse's microphone.

¹² No mice were trained on this task

¹³ (1920x1080px, 60Hz)

¹⁴ Visual stimuli were presented with Psychopy using the glfw backend while Autopilot was run in a dedicated X11 server.

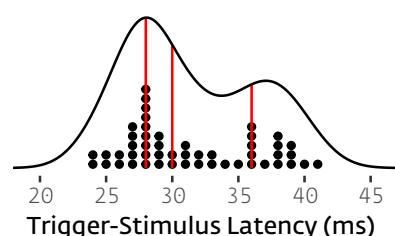


Figure 4.8: Stacked dots are a histogram of individual observations (n=50) underneath the probability density (black line), red lines indicate quartiles.

Autopilot is appropriate for realtime rendering of simple stimuli, and the proof-of-concept API we built around PsychoPy doesn't impose discernible overhead (Mean framerate for a 960 x 540px grating at 1080p in Autopilot: 36.2 fps, vs. timeByFrames: 35.0 fps). In the future we will investigate prerendering and caching complex stimuli in order to increase performance. A straightforward option for higher-performance video would be to deploy an Autopilot agent running on a desktop computer with a high-performance GPU, or to use a single-board computer with a GPU like the NVIDIA Jetson (\$99)¹⁵.

Stimulus	Resolution	Size / # Dots	Mean FPS	σ FPS
Gabor Filter	1280 x 720	300 x 300px	106.4	5.5
Gabor Filter	1920 x 1080	300 x 300px	75.2	3.5
Gabor Filter	1280 x 720	640 x 360px	53.5	2.2
Gabor Filter	1920 x 1080	960 x 540px	35.0	1.0
Gabor Filter	1280 x 720	720 x 720px	41.5	2.2
Gabor Filter	1920 x 1080	1080 x 1080px	20.1	0.7
Random Dots	1280 x 720	100 dots	98.0	3.8
Random Dots	1920 x 1080	100 dots	67.6	3.0
Random Dots	1280 x 720	1000 dots	20.9	0.25
Random Dots	1920 x 1080	1000 dots	19.5	0.36

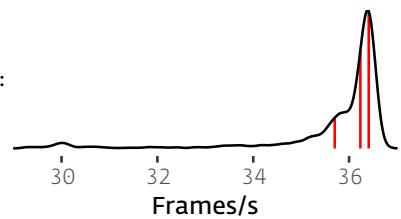


Figure 4.9: Probability density of framerates for 960 x 540px grating rendered at 1080p. Red lines indicate quartiles

¹⁵ as we did in [29]

Table 4.8: Tests performed over 1000 frames with PsychoPy's timeByFrames test.

1133 *Limitations and Future Directions*

1134 WE WILL LIKELY NEVER VIEW Autopilot as “finished.” Autopilot—like all open-source software—is an evolving project,
 1135 and this paper captures it as a snapshot at v0.5.0. We are invested in its development, and will be continually working to fix bugs,
 1136 make its use more elegant, and add new features in collaboration with other researchers.

1137 We expect that as the codebase matures and other researchers use Autopilot in new, unexpected ways that some fundamental
 1138 elements of its structure may evolve. We have built version logging into the structure of the system so that changes will not
 1139 compromise the replicability of experiments (see 5.5 below). While there will inevitably be breaking changes, these will be trans-
 1140 parently documented, announced in release notes, and indicated with semantic versioning in order to alert users and describe
 1141 how to adapt as needed.

1142 We recognize the risk and inertia of retooling lab infrastructure, and there is still much work to be done on Autopilot. We welcome
 1143 all issues and questions from anyone interested in contributing, or just curious to try it out — trying Autopilot is ultimately as
 1144 risky as buying a Raspberry Pi.

1145 The current major planned changes (also see the todo page in the docs) include:

- 1146 1. **Python, Meet Rust** - Python is very useful as a high-level glue language, and its accessibility to a large number of scientific
 1147 programmers is important to us, but it has its own very real performance limitations. As Autopilot’s modules mature and
 1148 stabilize, we are interested in rewriting core routines like sound presentation and networking in rust and exposing them to
 1149 python with tools like PyO3
- 1150 2. **Real Realtime** - Beneath user space decisions like programming language, the timing of CPU operations in linux is still
 1151 determined by the kernel — this is one of the major reasons why other projects are based around dedicated microcontrollers.
 1152 For almost everything that most scientists want to do, the standard linux kernel is perfectly fine, but we are interested in
 1153 investigating what it would take to provide true deterministic realtime performance via Autopilot’s high-level object system.
 1154 One approach might be to provide prebuilt realtime kernel images along with tools to easily deploy them, though no firm
 1155 plan has been made.
- 1156 3. **Integration** - We will continue to collaborate with other programming teams to be interoperable with a broader array of other
 1157 tools. Our next set of planned integrations include recording electrophysiological data by integrating with Open Ephys[57],
 1158 optical imaging data from the Miniscope project [58, 59], and shared processing and control pipelines with Bonsai[19].
- 1159 4. **Data Ingest & Export** - We are releasing Autopilot’s data modeling system in v0.5.0 as an alpha release alongside this paper,
 1160 and it includes prototype export interfaces to Neurodata Without Borders[32] and Datajoint[60]. Over the next several
 1161 releases, we will continue to improve our data model so that researchers can easily structure their data and choose among
 1162 different backends for storage. We are also working on a separate project to make tools to ingest data from the more ad-hoc
 1163 directory-based data formats widely used in science and ingest them into Autopilot’s and other tools formal modeling systems.
 1164 In the longer term, we are interested in making Autopilot interoperable with linked data systems as part of a broader vision
 1165 of digital infrastructure.
- 1166 5. **Provenance** - Autopilot stores version information and local configuration in multiple places, and it is technically possible
 1167 to faithfully replicate an experiment, but recording of provenance can still be consolidated and improved. By formalizing our
 1168 object and data model, we will also systematize the many changes in configuration and version possible across the system for
 1169 complete provenance tracking.
- 1170 6. **P2P Networking** - The default tree structure of Autopilot’s networking modules has proven to be unnecessarily limiting
 1171 over time. In part, we had preoptimized for processing messages in a separate processes assuming that would help problems
 1172 from dropped messages and overflowing send buffers, but in practice messages are almost never dropped and network nodes
 1173 are as effective as stations in sending and receiving large amounts of data. As part of unifying Autopilot’s object system, we will

1174 implement a fully peer-to-peer networking system such that each instantiated object has a unique ID so that messages can be
 1175 easily addressed from any object to any other in its swarm. We will learn from previous p2p addressing systems like [distributed](#)
 1176 [hash tables](#) to allow net nodes to join the swarm and discover all other nodes automatically without manually configuring IP
 1177 addresses and ports. In the longer term we are interested in peer to peer data transfer as well, so that an object serving as a data
 1178 source can efficiently stream to many consumers without needing to duplicate each message for every consumer.

- 1179 7. **Slots, Signals, and Streaming** - We will be supplementing a more general network structure with a system of specifying
 1180 which attributes of each object are data sources, which are sinks, and what kind of connection they accept. Similar to Qt's
 1181 [signal and slot](#) model, we want to make it as easy as using a `.connect()` method to control one piece of hardware with another.
 1182 The transforms module should also be able to support branching and forking operations so multiple data sources can be
 1183 combined for elaborated hardware control. ZeroMQ is an excellent tool for sending and receiving control messages, but
 1184 formalized signals and slots could also specify different streaming tools like [redis](#) or [gstreamer](#) that might be better suited for
 1185 high-bandwidth linear streams like video. Applied generally, this could also solve related problems like the relatively implicit
 1186 handling of event triggers in the Task class and the need for manual configuration of connections between pilots and copilots.
- 1187 8. **Rebuild the GUI** - The GUI is some of the oldest code in the library, was written before most of the other modules existed,
 1188 and needs to be rebuilt. We have started by remaking its central widgets to be [generated from pydantic models](#) used increasingly
 1189 throughout the system, but the rest of the GUI still needs to be rearchitected into a structure that decreases code duplication
 1190 and allows us to do things like provide GUI extensions via plugins. We will likely continue to use Qt for the near future, but
 1191 are also exploring the idea of webassembly tools to make browser-based web interfaces for remote control.
- 1192 9. **Plugins** - We want Autopilot's plugin system to be permissive and as natural as the scripting style that most experimental
 1193 code is written in, but we still need some means of specifying dependencies on other packages and plugins, among other
 1194 improvements. We will be making a plugin generator that makes a folder of plugin boilerplate, as well as tools for installing,
 1195 uploading, and synchronizing versions with git and the wiki. Over time we will make all object types within Autopilot able
 1196 to be extended with plugins, as well as make it possible to override and extend built-in objects.
- 1197 10. **Knowledge Organization** - We have been extending our thinking from code itself to more broadly consider the social systems
 1198 that surround research code. The wiki was our first step, and we will continue to make more points of integration for smoothly
 1199 incorporating contextual knowledge typically stored in lab notebooks into a public, collectively curated information system.
 1200 We want to make it easier not only for individual researchers to use Autopilot, but make it easier for labs to coordinate work
 1201 across projects without needing to rely on proprietary SaaS platforms with additional tooling for managing swarms, and
 1202 moving beyond a single Autopilot wiki to a federated system of wikis for fluid continuity between "private" local coordination
 1203 and "public" shared knowledge.
- 1204 11. **Tests** - Our collection of tests doesn't cover the whole codebase, and so as we formalize our contribution process will move
 1205 towards a system where all new code must have tests and documentation to be integrated. We also want to integrate our
 1206 tests more closely with our documentation so that researchers know which part of the code has explicit tests guaranteeing
 1207 functionality.
- 1208 12. **Security** - Autopilot is a networked program, and while it doesn't execute arbitrary code from network messages, there is
 1209 no security model to speak of. So far this hasn't been a problem, as we encourage only using Autopilot on a local network
 1210 behind a router, but as we build out our networking modules we will investigate how to incorporate identity verification
 1211 systems to protect swarms from malicious messages.
- 1212 13. **Metastructure & API Maturity** - The scope and structure of Autopilot is still in flux relative to other, more mature Python
 1213 packages. To reach a stable v1.0.0 API, we are in the process of unifying Autopilot's object structure so everything is clearly
 1214 typed, all configuration is explicit, and all code written to handle special cases is absorbed into more general systems. Different
 1215 parts of Autopilot have had different degrees of care over time, and so we will be working to catch the oldest modules up,
 1216 trim unused ones, and make sure every line in the library is documented and useful. For the time being, flexibility is useful
 1217 because frequently used or requested features trace a desire path outlining how its users believe Autopilot should behave. Each
 1218 shortcoming we fix in Autopilot's modules makes it more straightforward to fix the rest, and so once the major remaining
 1219 work is completed we will transition to a more conservative pace of development that ensures the longevity of the project.

Glossary

Agent	3.7	The executable part of Autopilot. A set of startup routines (eg. opening a GUI or starting an audio server), runtime behavior (eg. opening as a window or running as a background system process), and event handling methods (ie. listens) that constitute the role of the particular Autopilot instance in the swarm .
Copilot	3.7	An agent that performs some auxiliary, supporting role in a task —primarily used for offloading some hardware responsibilities from a pilot .
Graduation	3.3	Moving between successive tasks in a protocol when some criterion is met.
Listen	3.8	A method belonging to the station or node of a particular agent that defines how to process a particular type of message (ie. a message with a particular key).
Node	3.8	A networking object that some module (eg. hardware, tasks , GUI routines) or method (eg. a listen) uses to communicate with other nodes . Messages to other agents in the swarm are relayed through their Station
Pilot	3.7	An agent that runs on a Raspberry Pi, the primary experimental agent of Autopilot. Typically runs as a system service, receives tasks from a terminal and runs them. Can organize a group of children if requested by the task .
Protocol	3.3	A (.json) file that contains a list of task parameters and the graduation criteria to move between them. The tasks in a protocol are also known as its levels .
Stage	3.3	Stages are methods that implement the logic of a task . They can be used analogously to states in a finite-state machine (eg. wait for trial initiation, play stimulus, etc.) or asynchronously (whenever x input is received, rotate stimulus by y degrees).
Station	3.8	Each agent has a single station , a networking object that is run in its own process and is responsible for communication between agents . The station also routes messages from children or other nodes .
Swarm		Informally, a group of connected agents .
Task	3.3	A formalized description of an experiment: the parameters it takes, the data that it collects, the hardware it needs, and a collection of stages that describe what happens during the experiment.
Terminal	3.7	A user-facing agent that provides a GUI for operating and maintaining a swarm .
Topology	3.7	A particular combination of agents , their designated responsibilities, and the networking connections between them invoked by a task (eg. task requires one pilot to record video, one to process the video, and one to administer reward) or by usage (eg. 10 pilots are connected to a single terminal and are typically used to run 10 independent tasks, though they could run shared tasks together).
Trial	3.3	If a task is structured such that its stages form a repeating series, a trial is a single completion of that series.

1223 Bibliography

- 1224 [1] Philip Meier, Erik Flister, and Pamela Reinagel. Collinear features impair visual detection by rats. 11(3). ISSN 1534-7362.
1225 <https://doi.org/10.1167/11.3.22>. (document)
- 1226 [2] Santiago Jaramillo, Anna Lakunina, Lan Guo, and Nick Ponvert. TASKontrol. URL <https://github.com/sjara/taskkontrol>. (document)
- 1227
- 1228 [3] Jacob Reimer, Matthew J. McGinley, Yang Liu, Charles Rodenkirch, Qi Wang, David A. McCormick, and Andreas S. Tolias. Pupil fluctuations track rapid changes in adrenergic and cholinergic activity in cortex. 7:13289. ISSN 2041-1723.
1229 <https://doi.org/10.1038/ncomms13289>. URL <https://www.nature.com/articles/ncomms13289>. 1
- 1230
- 1231 [4] Ana Parabucki, Alexander Bizer, Genela Morris, Antonio E. Munoz, Avinash D. S. Bala, Matthew Smear, and Roman Shusterman. Odor Concentration Change Coding in the Olfactory Bulb. 6(1). ISSN 2373-2822. <https://doi.org/10.1523/JNEURO.0396-18.2019>. URL <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC6397952/>. 1
- 1232
- 1233
- 1234 [5] Christopher M. Niell and Michael P. Stryker. Modulation of Visual Responses by Behavioral State in Mouse Visual Cortex. 65(4):472–479. ISSN 0896-6273. <https://doi.org/10.1016/j.neuron.2010.01.033>. URL <http://www.sciencedirect.com/science/article/pii/S0896627310000590>. 1
- 1235
- 1236
- 1237 [6] Tanmay Nath, Alexander Mathis, An Chi Chen, Amir Patel, Matthias Bethge, and Mackenzie Weygandt Mathis. Using DeepLabCut for 3D markerless pose estimation across species and behaviors. 14(7):2152–2176. ISSN 1750-2799. <https://doi.org/10.1038/s41596-019-0176-0>. URL <https://www.nature.com/articles/s41596-019-0176-0>. 1
- 1238
- 1239
- 1240 [7] James J. Jun, Nicholas A. Steinmetz, Joshua H. Siegle, Daniel J. Denman, Marius Bauza, Brian Barbarits, Albert K. Lee, Costas A. Anastassiou, Alexandru Andrei, Çağatay Aydin, Mladen Barbic, Timothy J. Blanche, Vincent Bonin, João Couto, Barundeb Dutta, Sergey L. Gratiy, Diego A. Gutnisky, Michael Häusser, Bill Karsh, Peter Ledochowitsch, Carolina Mora Lopez, Catalin Mitelut, Silke Musa, Michael Okun, Marius Pachitariu, Jan Putzeys, P. Dylan Rich, Cyrille Rossant, Wei-Lung Sun, Karel Svoboda, Matteo Carandini, Kenneth D. Harris, Christof Koch, John O’Keefe, and Timothy D. Harris. Fully integrated silicon probes for high-density recording of neural activity. 551(7679):232–236. ISSN 1476-4687. <https://doi.org/10.1038/nature24636>. 1, 2, 5, 11
- 1241
- 1242
- 1243
- 1244
- 1245
- 1246
- 1247 [8] Christopher P. Burgess, Armin Lak, Nicholas A. Steinmetz, Peter Zatka-Haas, Charu Bai Reddy, Elina A. K. Jacobs, Jennifer F. Linden, Joseph J. Paton, Adam Ranson, Sylvia Schröder, Sofia Soares, Miles J. Wells, Lauren E. Wool, Kenneth D. Harris, and Matteo Carandini. High-Yield Methods for Accurate Two-Alternative Visual Psychophysics in Head-Fixed Mice. 20(10):2513–2524. ISSN 2211-1247. <https://doi.org/10.1016/j.celrep.2017.08.047>. URL <http://www.sciencedirect.com/science/article/pii/S2211124717311725>. 1
- 1248
- 1249
- 1250
- 1251
- 1252 [9] Kay Thurley and Aslı Ayaz. Virtual reality systems for rodents. 63(1):109–119. ISSN 1674-5507. <https://doi.org/10.1093/cz/zow070>. URL <https://academic.oup.com/cz/article/63/1/109/2962415>. 1
- 1253
- 1254 [10] Anna R. Chambers, Kenneth E. Hancock, Kamal Sen, and Daniel B. Polley. Online stimulus optimization rapidly reveals multidimensional selectivity in auditory cortical neurons. 34(27):8963–8975. ISSN 1529-2401. <https://doi.org/10.1523/JNEUROSCI.0260-14.2014>. 1
- 1255
- 1256
- 1257 [11] Chance Elliott, Vipin Vijayakumar, Wesley Zink, and Richard Hansen. National Instruments LabVIEW: A Programming Environment for Laboratory Automation and Measurement. 12(1):17–24. ISSN 1535-5535. <https://doi.org/10.1016/j.jala.2006.07.012>. URL <https://journals.sagepub.com/doi/abs/10.1016/j.jala.2006.07.012>. 1
- 1258
- 1259
- 1260 [12] Open Ephys. pyControl. URL <http://www.open-ephys.org/store/pycontrol>. 1
- 1261 [13] Josh Sanders. Sanworks - BPod. URL <https://www.sanworks.io/shop/products.php?productFamily=bpod>. 1
- 1262

- 1262 [14] Matthew B. Wall. Reliability starts with the experimental tools employed. 113:352–354. ISSN 0010-9452.
 1263 <https://doi.org/10.1016/j.cortex.2018.11.034>. URL <http://www.sciencedirect.com/science/article/pii/S001094521830443X>. 1, 2.2
- 1265 [15] Peter Johnson-Lenz and Trudy Johnson-Lenz. Post-mechanistic groupware primitives: Rhythms, boundaries and contain-
 1266 ers. 34(3):395–417,. ISSN 0020-7373. [https://doi.org/10.1016/0020-7373\(91\)90027-5](https://doi.org/10.1016/0020-7373(91)90027-5). URL [https://doi.org/10.1016/0020-7373\(91\)90027-5](https://doi.org/10.1016/0020-7373(91)90027-5). 1
- 1268 [16] Peter Johnson-Lenz and Trudy Johnson-Lenz. Groupware: Coining and defining it. 19(2):34,. ISSN 2372-7403, 2372-
 1269 739X. <https://doi.org/10.1145/290575.290585>. URL <https://dl.acm.org/doi/10.1145/290575.290585>. 1
- 1270 [17] STEPHEN R. BARLEY and BETH A. BECHKY. In the Backrooms of Science: The Work of Technicians in Science
 1271 Labs. 21(1):85–126. ISSN 0730-8884. <https://doi.org/10.1177/0730888494021001004>. URL <https://doi.org/10.1177/0730888494021001004>. 1
- 1273 [18] Thomas Akam, Andy Lustig, James M Rowland, Sampath KT Kapanaiah, Joan Esteve-Agraz, Mariangela Panniello,
 1274 Cristina Márquez, Michael M Kohl, Dennis Kätsel, Rui M Costa, and Mark E Walton. Open-source, Python-based,
 1275 hardware and software for controlling behavioural neuroscience experiments. 11:e67846. ISSN 2050-084X. <https://doi.org/10.7554/eLife.67846>. URL <https://doi.org/10.7554/eLife.67846>. 1.1
- 1277 [19] Gonçalo Lopes, Niccolò Bonacchi, João Frazão, Joana P. Neto, Bassam V. Atallah, Sofia Soares, Luís Moreira, Sara Matias,
 1278 Pavel M. Itskov, Patrícia A. Correia, Roberto E. Medina, Lorenza Calcaterra, Elena Dreosti, Joseph J. Paton, and Adam R.
 1279 Kampff. Bonsai: An event-based framework for processing and controlling data streams. 9. ISSN 1662-5196. URL
 1280 <https://www.frontiersin.org/article/10.3389/fninf.2015.00007>. 5, 5.3
- 1281 [20] Florian Krause and Oliver Lindemann. Expyriment: A Python library for cognitive and neuroscientific experiments. 46
 1282 (2):416–428. ISSN 1554-3528. <https://doi.org/10.3758/s13428-013-0390-6>. URL <https://doi.org/10.3758/s13428-013-0390-6>. 5
- 1284 [21] Jonathan Peirce, Jeremy R. Gray, Sol Simpson, Michael MacAskill, Richard Höchenberger, Hiroyuki Sogo, and Erik Kast-
 1285 man. PsychoPy2: Experiments in behavior made easy. 51(1):195–203. ISSN 1554-3528. <https://doi.org/10.3758/s13428-018-01193-y>. URL <https://link.springer.com/article/10.3758/s13428-018-01193-y>. 5, 3.6
- 1287 [22] Sebastiaan Mathôt, Daniel Schreij, and Jan Theeuwes. OpenSesame: An open-source, graphical experiment builder for
 1288 the social sciences. 44(2):314–324. ISSN 1554-3528. <https://doi.org/10.3758/s13428-011-0168-7>. URL <https://doi.org/10.3758/s13428-011-0168-7>. 5
- 1290 [23] Xinfeng Chen and Haohong Li. ArControl: An Arduino-Based Comprehensive Behavioral Platform with Real-Time
 1291 Performance. 11. ISSN 1662-5153. <https://doi.org/10.3389/fnbeh.2017.00244>. URL <https://www.frontiersin.org/articles/10.3389/fnbeh.2017.00244/full>. 5
- 1293 [24] S. van der Walt, S. C. Colbert, and G. Varoquaux. The NumPy Array: A Structure for Efficient Numerical Computation.
 1294 13(2):22–30. ISSN 1521-9615. <https://doi.org/10.1109/MCSE.2011.37>. 6, 4.1
- 1295 [25] Eric Jones, Travis Oliphant, and Pearu Peterson. SciPy: Open Source Scientific Tools for Python. URL <http://www.scipy.org/>. 6
- 1297 [26] Sandeep Robert Datta, David J. Anderson, Kristin Branson, Pietro Perona, and Andrew Leifer. Computational Neu-
 1298 roethology: A Call to Action. 104(1):11–24. ISSN 0896-6273. <https://doi.org/10.1016/j.neuron.2019.09.038>.
 1299 URL [https://www.cell.com/neuron/abstract/S0896-6273\(19\)30841-4](https://www.cell.com/neuron/abstract/S0896-6273(19)30841-4). 1.2, 3.7
- 1300 [27] Sanworks, LLC. 8 reasons to use Bpod’s new HiFi module. URL <https://sanworks.io/news/viewArticle.php?articleID=HiFi01>. 1.2

- [28] The International Brain Laboratory, Valeria Aguillon-Rodriguez, Dora E. Angelaki, Hannah M. Bayer, Niccolò Bonacchi, Matteo Carandini, Fanny Cazettes, Gaelle A. Chapuis, Anne K. Churchland, Yang Dan, Eric E. J. Dewitt, Mayo Faulkner, Hamish Forrest, Laura M. Haetzel, Michael Hausser, Sonja B. Hofer, Fei Hu, Anup Khanal, Christopher S. Krasniak, Inês Laranjeira, Zachary F. Mainen, Guido T. Meijer, Nathaniel J. Miska, Thomas D. Mrsic-Flogel, Masayoshi Murakami, Jean-Paul Noel, Alejandro Pan-Vazquez, Cyrille Rossant, Joshua I. Sanders, Karolina Z. Socha, Rebecca Terry, Anne E. Urai, Hernando M. Vergara, Miles J. Wells, Christian J. Wilson, Ilana B. Witten, Lauren E. Wool, and Anthony Zador. Standardized and reproducible measurement of decision-making in mice. page 2020.01.17.909838. <https://doi.org/10.1101/2020.01.17.909838>. URL <https://www.biorxiv.org/content/10.1101/2020.01.17.909838v5>. 1.2
- [29] Gary A Kane, Gonçalo Lopes, Jonny L Saunders, Alexander Mathis, and Mackenzie W Mathis. Real-time, low-latency closed-loop feedback using markerless posture tracking. 9:e61909. ISSN 2050-084X. <https://doi.org/10.7554/eLife.61909>. URL <https://doi.org/10.7554/eLife.61909>. 2.1, 3.5, 3.5, 3.6, 15
- [30] Pieter Hintjens. *ZeroMQ: Messaging for Many Applications*. O'Reilly Media, 1 edition edition. ISBN 978-1-4493-3406-2. 2.1, 11
- [31] Jonny L. Saunders and Michael Wehr. Mice can learn phonetic categories. 145(3):1168–1177. ISSN 0001-4966. <https://doi.org/10.1121/1.5091776>. URL <https://asa.scitation.org/doi/abs/10.1121/1.5091776>. 2.1
- [32] Oliver Rübel, Andrew Tritt, Benjamin Dichter, Thomas Braun, Nicholas Cain, Nathan Clack, Thomas J. Davidson, Max Dougherty, Jean-Christophe Fillion-Robin, Nile Graddis, Michael Grauer, Justin T. Kiggins, Lawrence Niu, Doruk Ozturk, William Schroeder, Ivan Soltesz, Friedrich T. Sommer, Karel Svoboda, Ng Lydia, Loren M. Frank, and Kristofer Bouchard. NWB:N 2.0: An Accessible Data Standard for Neurophysiology. <https://doi.org/10.1101/523035>. URL <https://www.biorxiv.org/content/10.1101/523035v1>. 2.2, 5.4
- [33] David A. W. Soergel. Rampant software errors may undermine scientific results. 3. ISSN 2046-1402. <https://doi.org/10.12688/f1000research.5930.2>. URL <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4629271/>. 2.2
- [34] Anders Eklund, Thomas E. Nichols, and Hans Knutsson. Cluster failure: Why fMRI inferences for spatial extent have inflated false-positive rates. 113(28):7900–7905. ISSN 0027-8424, 1091-6490. <https://doi.org/10.1073/pnas.1602413113>. URL <https://www.pnas.org/content/113/28/7900>. 2.2
- [35] Jayanti Bhandari Neupane, Ram P. Neupane, Yuheng Luo, Wesley Y. Yoshida, Rui Sun, and Philip G. Williams. Characterization of Leptazolines A–D, Polar Oxazolines from the Cyanobacterium Leptolyngbya sp., Reveals a Glitch with the “Willoughby–Hoye” Scripts for Calculating NMR Chemical Shifts. 21(20):8449–8453. ISSN 1523-7060. <https://doi.org/10.1021/acs.orglett.9b03216>. URL <https://doi.org/10.1021/acs.orglett.9b03216>. 2.2
- [36] Greg Miller. A Scientist’s Nightmare: Software Problem Leads to Five Retractions. 314(5807):1856–1857. ISSN 0036-8075, 1095-9203. <https://doi.org/10.1126/science.314.5807.1856>. URL <https://science.sciencemag.org/content/314/5807/1856>. 2.2
- [37] Yarden Katz and Ulrich Bernhard Matter. On the Biomedical Elite: Inequality and Stasis in Scientific Knowledge Production. URL <http://nrs.harvard.edu/urn-3:HUL.InstRepos:33373356>. 2.2
- [38] Jeremy Ashkenas, Haeyoun Park, and Adam Pearce. Even With Affirmative Action, Blacks and Hispanics Are More Underrepresented at Top Colleges Than 35 Years Ago. ISSN 0362-4331. URL <https://www.nytimes.com/interactive/2017/08/24/us/affirmative-action.html>, <https://www.nytimes.com/interactive/2017/08/24/us/affirmative-action.html>. 2.2
- [39] Aaron Clauset, Samuel Arbesman, and Daniel B. Larremore. Systematic inequality and hierarchy in faculty hiring networks. 1(1):e1400005. ISSN 2375-2548. <https://doi.org/10.1126/sciadv.1400005>. URL <https://advances.sciencemag.org/content/1/1/e1400005>. 2.2
- [40] J. M. Pearce, J. C. Molloy, S. Kuznetsov, and S. Dosemagen. Expanding Equitable Access to Experimental Research and STEM Education by Supporting Open

- 1345 Source Hardware Development. URL <http://openhardware.science/2019/01/27/expanding-equitable-access-to-experimental-research-and-stem-education-by-supporting-open-source-hardware-development-2.2>
- 1348 [41] Emmeke Aarts, Matthijs Verhage, Jesse V. Veenvliet, Conor V. Dolan, and Sophie van der Sluis. A solution to dependency: Using multilevel analysis to accommodate nested data. 17(4):491–496. ISSN 1546-1726. <https://doi.org/10.1038/nrn3648>. URL <https://www.nature.com/articles/nrn3648>. 2.5, 2.2
- 1351 [42] Patrick E. Shrout and Joseph L. Rodgers. Psychology, Science, and Knowledge Construction: Broadening Perspectives from the Replication Crisis. 69(1):487–510. <https://doi.org/10.1146/annurev-psych-122216-011845>. URL <https://doi.org/10.1146/annurev-psych-122216-011845>. 2.2
- 1354 [43] Katherine S. Button, John P. A. Ioannidis, Claire Mokrysz, Brian A. Nosek, Jonathan Flint, Emma S. J. Robinson, and Marcus R. Munafò. Power failure: Why small sample size undermines the reliability of neuroscience. 14(5):365–376. ISSN 1471-0048. <https://doi.org/10.1038/nrn3475>. URL <https://www.nature.com/articles/nrn3475>. 2.2
- 1357 [44] Tim Anderson. Guido van Rossum aiming to make CPython 2x faster in 3.11. URL https://www.theregister.com/2021/05/13/guido_van_rossum_cpython_3_11/. 2
- 1359 [45] Wes McKinney. Pandas: A foundational Python library for data analysis and statistics. 14(9):1–9. URL https://www.dlr.de/sc/portaldatas/15/resources/dokumente/pyhpc2011/submissions/pyhpc2011_submission_9.pdf. 3.2, 4.1
- 1361 [46] Dexter C. Kozen. Limitations of Finite Automata. In Dexter C. Kozen, editor, Automata and Computability, Undergraduate Texts in Computer Science, pages 67–71. Springer New York. ISBN 978-1-4612-1844-9. https://doi.org/10.1007/978-1-4612-1844-9_12. 3.3
- 1364 [47] Ji Hyun Bak, Jung Yoon Choi, Athena Akrami, Ilana Witten, and Jonathan W Pillow. Adaptive optimal training of animal behavior. In D. D. Lee, M. Sugiyama, U. V. Luxburg, I. Guyon, and R. Garnett, editors, Advances in Neural Information Processing Systems 29, pages 1947–1955. Curran Associates, Inc. URL <http://papers.nips.cc/paper/6344-adaptive-optimal-training-of-animal-behavior.pdf>. 3.3
- 1368 [48] Aaron Swartz. Aaron Swartz’s A Programmable Web: An Unfinished Work. 3(2):1–64. ISSN 2160-4711, 2160-472X. <https://doi.org/10.2200/S00481ED1V01Y201302WBE005>. URL <http://www.morganclaypool.com/doi/abs/10.2200/S00481ED1V01Y201302WBE005>. 6
- 1371 [49] Fatemeh Abyarjoo, Armando Barreto, Jonathan Cofino, and Francisco R. Ortega. Implementing a Sensor Fusion Algorithm for 3D Orientation Detection with Inertial/Magnetic Sensors. pages 305–310. https://doi.org/10.1007/978-3-319-06773-5_41. URL https://link.springer.com/chapter/10.1007/978-3-319-06773-5_41. 3.5
- 1374 [50] Photis Patonis, Petros Patias, Ilias N. Tziavos, Dimitrios Rossikopoulos, and Konstantinos G. Margaritis. A Fusion Method for Combining Low-Cost IMU/Magnetometer Outputs for Use in Applications on Mobile Devices. 18(8). ISSN 1424-8220. <https://doi.org/10.3390/s18082616>. URL <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC6111698/>. 3.5
- 1377 [51] Leland Wilkinson. The Grammar of Graphics. In James E. Gentle, Wolfgang Karl Härdle, and Yuichi Mori, editors, Handbook of Computational Statistics: Concepts and Methods, Springer Handbooks of Computational Statistics, pages 375–414. Springer Berlin Heidelberg. ISBN 978-3-642-21551-3. URL https://doi.org/10.1007/978-3-642-21551-3_13. 3.9
- 1381 [52] Hadley Wickham. Ggplot2: Elegant Graphics for Data Analysis. Springer-Verlag New York. ISBN 978-3-319-24277-4. URL <https://ggplot2.tidyverse.org>. 4.1
- 1383 [53] Hadley Wickham, Romain François, Lionel Henry, and Kirill Müller. Dplyr: A Grammar of Data Manipulation. URL <https://CRAN.R-project.org/package=dplyr>. 4.1
- 1385 [54] Lionel Henry and Hadley Wickham. Purrr: Functional Programming Tools. URL <https://CRAN.R-project.org/package=purrr>. 4.1

- 1387 [55] Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric
 1388 Wieser, Julian Taylor, Sebastian Berg, Nathaniel J. Smith, Robert Kern, Matti Picus, Stephan Hoyer, Marten H. van Kerk-
 1389 wijk, Matthew Brett, Allan Haldane, Jaime Fernández del Río, Mark Wiebe, Pearu Peterson, Pierre Gérard-Marchant,
 1390 Kevin Sheppard, Tyler Reddy, Warren Weckesser, Hameer Abbasi, Christoph Gohlke, and Travis E. Oliphant. Array pro-
 1391 gramming with NumPy. 585(7825):357–362. ISSN 1476-4687. <https://doi.org/10.1038/s41586-020-2649-2>. URL
 1392 <https://www.nature.com/articles/s41586-020-2649-2>. 4.1
- 1393 [56] Eduardo Soares, Pedro Brandão, and Rui Prior. Analysis of Timekeeping in Experimentation. In 2020 12th
 1394 International Symposium on Communication Systems, Networks and Digital Signal Processing (CSNDSP), pages 1–6.
 1395 <https://doi.org/10.1109/CSNDSP49049.2020.9249632>. 4.3, 3
- 1396 [57] Joshua H. Siegle, Aarón Cuevas López, Yogi A. Patel, Kirill Abramov, Shay Ohayon, and Jakob Voigts. OpenEphys:
 1397 An open-source, plugin-based platform for multichannel electrophysiology. 14(4):045003. ISSN 1741-2552. <https://doi.org/10.1088/1741-2552/aa5eea>. URL <https://doi.org/10.1088/1741-2552/aa5eea>. 5.3
- 1398 [58] Daniel Aharoni and Tycho M. Hoogland. Circuit Investigations With Open-Source Miniaturized Microscopes: Past,
 1399 Present and Future. 13. ISSN 1662-5102. URL <https://www.frontiersin.org/article/10.3389/fncel.2019.00141>.
 1400 5.3
- 1401 [59] Daniel Aharoni, Baljit S. Khakh, Alcino J. Silva, and Peyman Golshani. All the light that we can see: A new era in
 1402 miniaturized microscopy. 16(1):11–13. ISSN 1548-7105. <https://doi.org/10.1038/s41592-018-0266-x>. URL
 1403 <https://www.nature.com/articles/s41592-018-0266-x>. 5.3
- 1404 [60] Dimitri Yatsenko, Edgar Y. Walker, and Andreas S. Tolias. DataJoint: A Simpler Relational Data Model. URL <http://arxiv.org/abs/1807.11104>. 5.4
- 1405
- 1406